

Fundamentals of Machine Learning

Concepts, Techniques and Tools to Build Intelligent Systems

Module 5

Logistic Regression, Decision Tree &
Supported Vector Machine(SVM)

Ali Samanipour

May. 2023

1

What is an Imbalanced Dataset?

2

Normalization

3

Scaling

4

Splitting Dataset

5

Handling Imbalance

6

Logistic Regression

7

Decision Tree

8

Support Vector Machine (SVM)

9

Data Preparation & Metric Trap

10

Training & Evaluation

Credit Card Fraud Detection

The dataset has been collected and analyzed during a research collaboration of Worldline and the Machine Learning Group (<http://mlg.ulb.ac.be>) of ULB (Université Libre de Bruxelles) on big data mining and fraud detection

```
import pandas as pd
df = pd.read_csv("input/creditcard.csv")
df.shape
```

```
(284807, 31)
```

Let's Understand the Data

Time: Number of seconds elapsed between this transaction and the first transaction in the dataset

```
df.columns
```

```
Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',  
      'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',  
      'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',  
      'Class'],  
      dtype='object')
```

Let's Understand the Data ...

Amount: Transaction amount

```
df.columns
```

```
Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',  
      'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',  
      'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',  
      'Class'],  
      dtype='object')
```

Let's Understand the Data ...

Class: 1 for fraudulent transactions, 0 non-fraudulent

```
df.columns
```

```
Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',  
      'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',  
      'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',  
      'Class'],  
      dtype='object')
```

Let's Understand the Data ...

V1..., V28: These are the features of the dataset
(we will only be dependent on statistical metrics and relations to choose the best features for predicting the desired class)

```
df.columns
```

```
Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',  
      'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',  
      'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',  
      'Class'],  
      dtype='object')
```


1

What is an Imbalanced Dataset?

2

Normalization

3

Scaling

4

Splitting Dataset

5

Handling Imbalance

What is an Imbalanced Dataset?

An imbalanced dataset is when one output class has extremely high entries compared to the other output class

```
df.Class.value_counts()
```

```
0      284315
```


```
1       492
```

```
Name: Class, dtype: int64
```

Imbalanced Dataset

Let's see how
Imbalance Dataset
looks like

```
sns.countplot('Class', data=df)  
plt.title('Class Distributions \n (0: No Fraud || 1: Fraud)', fontsize=14)  
Text(0.5, 1.0, 'Class Distributions \n (0: No Fraud || 1: Fraud)')
```



Counting Nulls

We do not have any NULL values to handle in the dataset. So, going ahead

```
df.isnull().sum().max()
```

0

Knowing the Features

Let us see some of the values for all the features and know, how the value looks

```
df.sort_index(axis=1).head(3)
```

	Amount	Class	Time	V1	V10	V11	V12	V13	V14	V15	...	V26
0	149.62	0	0.0	-1.359807	0.090794	-0.551600	-0.617801	-0.991390	-0.311169	1.468177	...	-0.189115
1	2.69	0	0.0	1.191857	-0.166974	1.612727	1.065235	0.489095	-0.143772	0.635558	...	0.125895
2	378.66	0	1.0	-1.358354	0.207643	0.624501	0.066084	0.717293	-0.165946	2.345865	...	-0.139097

3 rows x 31 columns

Dataset Information

It's a relatively **large dataset**.

Training a model with all the features generally don't make sense.


We need to choose the features in such a way that those features will have some contribution to the decision of the output class.

```
df.sort_index(axis=1).info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
Amount      284807 non-null float64
Class       284807 non-null int64
Time        284807 non-null float64
V1          284807 non-null float64
V10         284807 non-null float64
V11         284807 non-null float64
V12         284807 non-null float64
V13         284807 non-null float64
V14         284807 non-null float64
V15         284807 non-null float64
V16         284807 non-null float64
```


Correlation of heat map

We now need to find some features among 28 dimensionally reduced features that we can use to train the model



Correlation of heat map ...

Features like V1, V3,
V5, V7, V9, V10,
V12, V14, V16, V17,
V18 have some
relation with
“Class” compared
to the other
features



1

What is an Imbalanced Dataset?

2

Normalization

3

Scaling

4

Splitting Dataset

5

Handling Imbalance

Normalization & Feature Scaling

All the 28 features are dimensionally reduced features through Principal Component Analysis (PCA). Before PCA is applied, the dataset is first normalized. So, we need to apply normalization to rest of the columns as well

```
df.sort_index(axis=1).head(3)
```

	Amount	Class	Time	V1	V10	V11	V12	V13	V14	V15	...	V26
0	149.62	0	0.0	-1.359807	0.090794	-0.551600	-0.617801	-0.991390	-0.311169	1.468177	...	-0.189115
1	2.69	0	0.0	1.191857	-0.166974	1.612727	1.065235	0.489095	-0.143772	0.635558	...	0.125895
2	378.66	0	1.0	-1.358354	0.207643	0.624501	0.066084	0.717293	-0.165946	2.345865	...	-0.139097

3 rows x 31 columns

Normalization & Feature Scaling

Normalization is used in a variety of ways in statistics. It is also known as **feature scaling**. The meaning we are going after, is the normalization of the range of the data to a standard scale.

Case 1: Say we have four lengths; all of them are in centimeters, and only one of them is inches. So, in those cases, we need to follow one **standard and generalize all** the length **units**.

Case 2: We have two features like Age and Salary, and we can easily say that both of the features will have a different range. Hence, we will use normalization/scaling techniques **to bring them on the same scale**

Range for all columns

Seeing the feature ranges, we can see “Time” and “Amount” has not been scaled.

df.min()		df.max()	
Time	0.000000	Time	172792.000000
V1	-56.407510	V1	2.454930
V2	-72.715728	V2	22.057729
V3	-48.325589	V3	9.382558
V4	-5.683171	V4	16.875344
V5	-113.743307	V5	34.801666
V6	-26.160506	V6	73.301626
V7	-43.557242	V7	120.589494
V8	-73.216718	V8	20.007208
V9	-13.434066	V9	15.594995
V10	-24.588262	V10	23.745136
V11	-4.797473	V11	12.018913
V12	-18.683715	V12	7.848392
V13	-5.791881	V13	7.126883
V14	-19.214325	V14	10.526766
V15	-4.498945	V15	8.877742
V16	-14.129855	V16	17.315112
V17	-25.162799	V17	9.253526
V18	-9.498746	V18	5.041069
V19	-7.213527	V19	5.591971
V20	-54.497720	V20	39.420904
V21	-34.830382	V21	27.202839
V22	-10.933144	V22	10.503090
V23	-44.807735	V23	22.528412
V24	-2.836627	V24	4.584549
V25	-10.295397	V25	7.519589
V26	-2.604551	V26	3.517346
V27	-22.565679	V27	31.612198
V28	-15.430084	V28	33.847808
Amount	0.000000	Amount	25691.160000
Class	0.000000	Class	1.000000
dtype: float64		dtype: float64	

Min-max Feature Scaling (Rescaling)

Here we scale the entire column with any given range into a usable given range $[a, b]$.

$$x' = a + \frac{(x_i - \min(x))(b - a)}{\max(x) - \min(x)}$$

Scaling down using min-max feature scaling, won't change the distribution of the feature. It only changes the range and keeping the distribution the same for the feature

Standardization (Z-score Normalization)

In the process of standardization, each feature has **zero-mean** and **unit variance or standard deviation**

$$x' = \frac{x_i - \bar{x}}{\sigma}$$

Scaling down using standardization will change the distribution of the feature and try to make mean close to zero, and the standard deviation equals to one

Principal Component Analysis(PCA)

Principal component analysis (PCA) is a statistical method to explain variance and covariance structure of a set of variables through linear combination. It uses the concept of orthogonal transformation to convert the set of variables (mostly correlated variables) into a set of linearly uncorrelated values

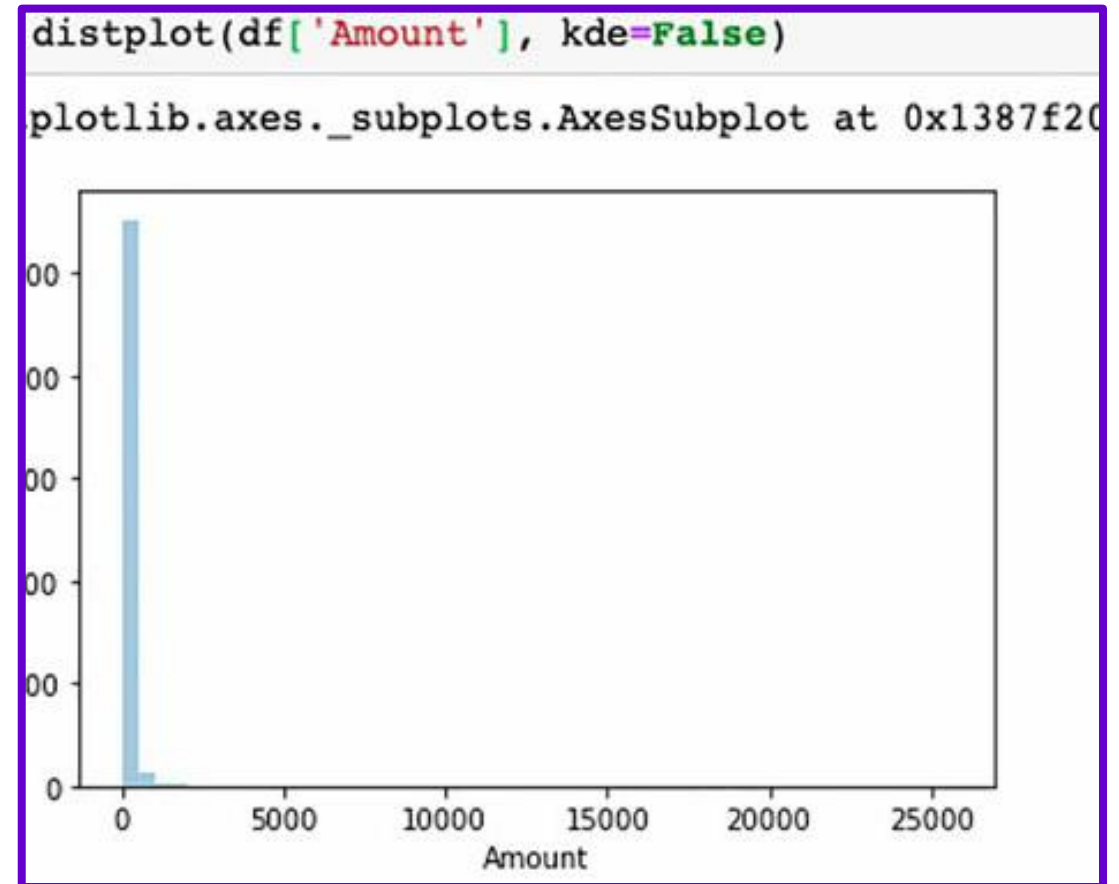
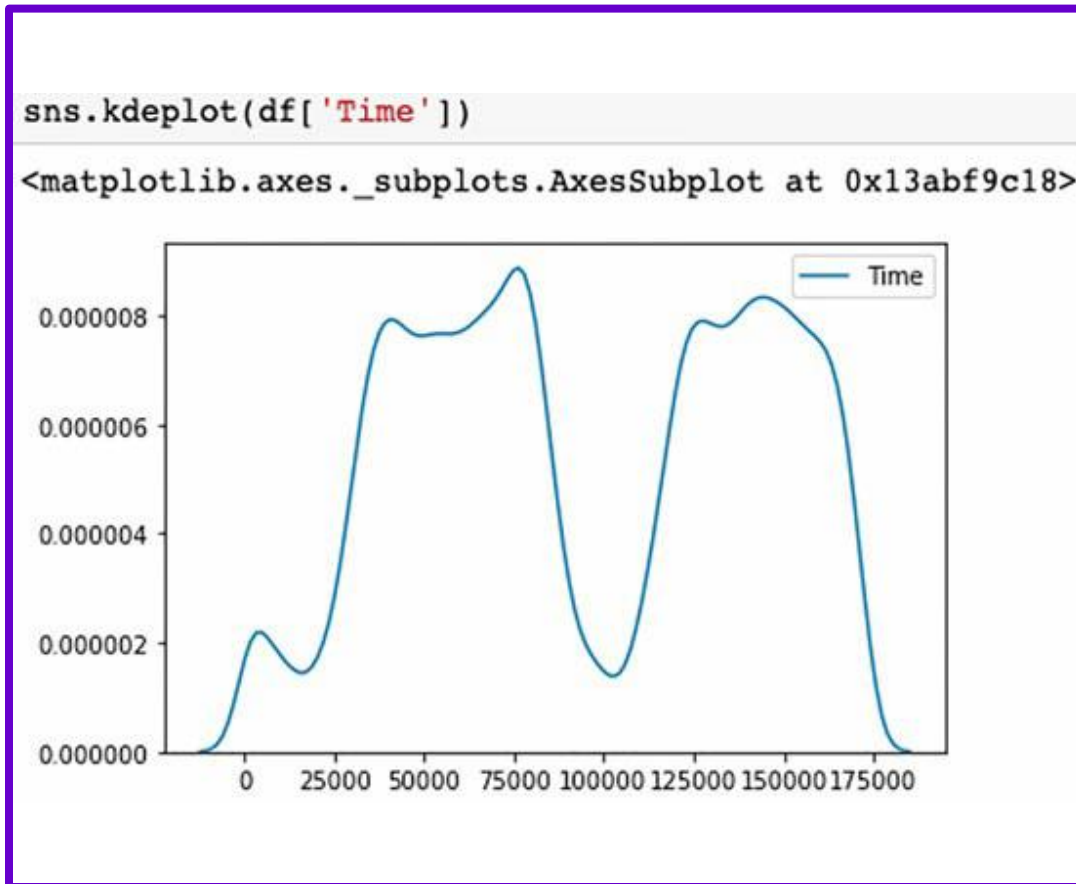
PCA is one of the popular **methods to perform dimension reduction on a large dataset**, and it helps in multiple ways like visualization, handling a smaller number of columns/features for modeling

Cross-validation



Cross-validation is a model validation technique to **ensure its performance**, but it only works best when the training dataset distribution is similar to real-time data distribution.

There are different types of cross-validation techniques: Exhaustive cross-validation like Leave-one-out cross-validation and Non-exhaustive cross-validation like k-fold cross-validation

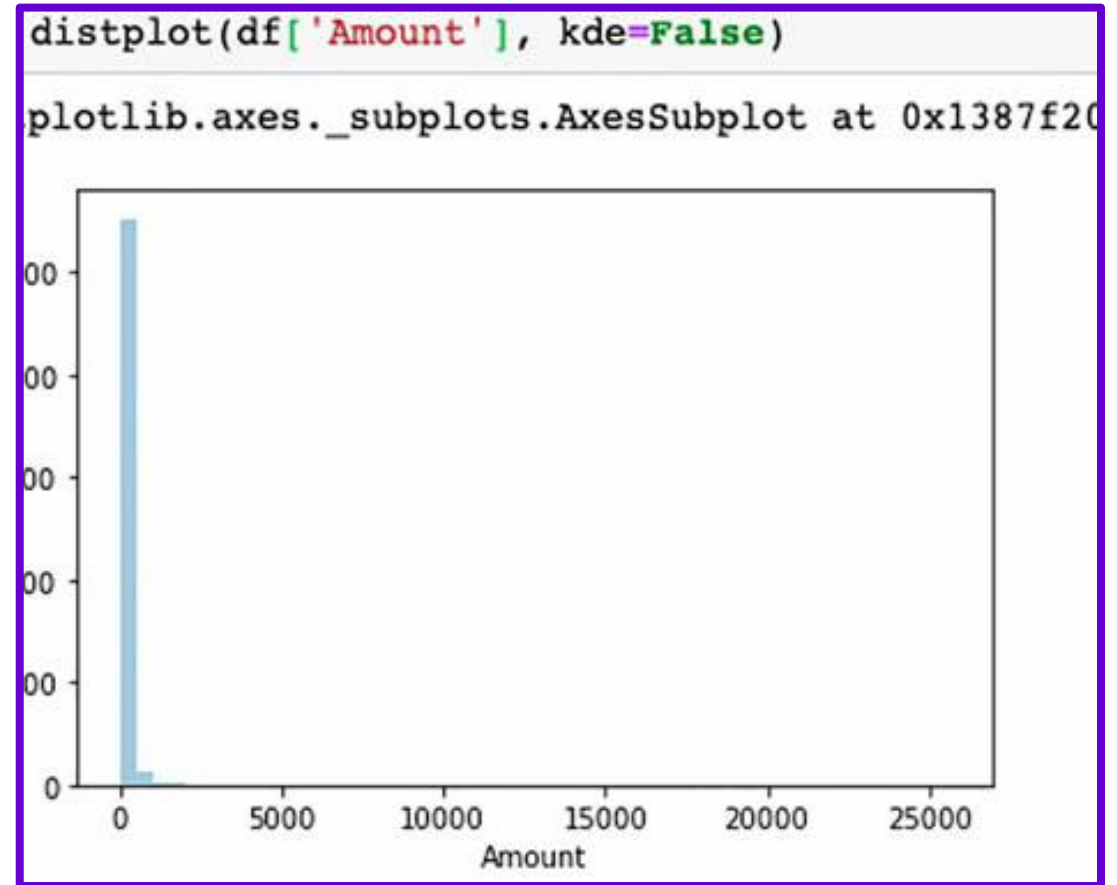

Bi-modal “Time” distribution plot



“Amount” Frequency Distribution




“Amount” Frequency Distribution



Comparison for an “Amount” and “Time” KDEs

We should try to
see if there is any
visual link/similarity
among the
distributions.

```
fig, ax = plt.subplots(1, 2, figsize=(14,4))  
  
sns.distplot(df['Amount'], ax=ax[0], color='r')  
ax[0].set_title('Distribution of Transaction Amount', fontsize=14)  
  
sns.distplot(df['Time'], ax=ax[1], color='b')  
ax[1].set_title('Distribution of Transaction Time', fontsize=14)  
  
plt.show()
```



Summed “Amount”
grouped by “Time

we need to drop
this idea to
compare “Time”
and “Amount”
distribution
separately.

```
df[['Time', 'Amount']].groupby(['Time']).head()
```

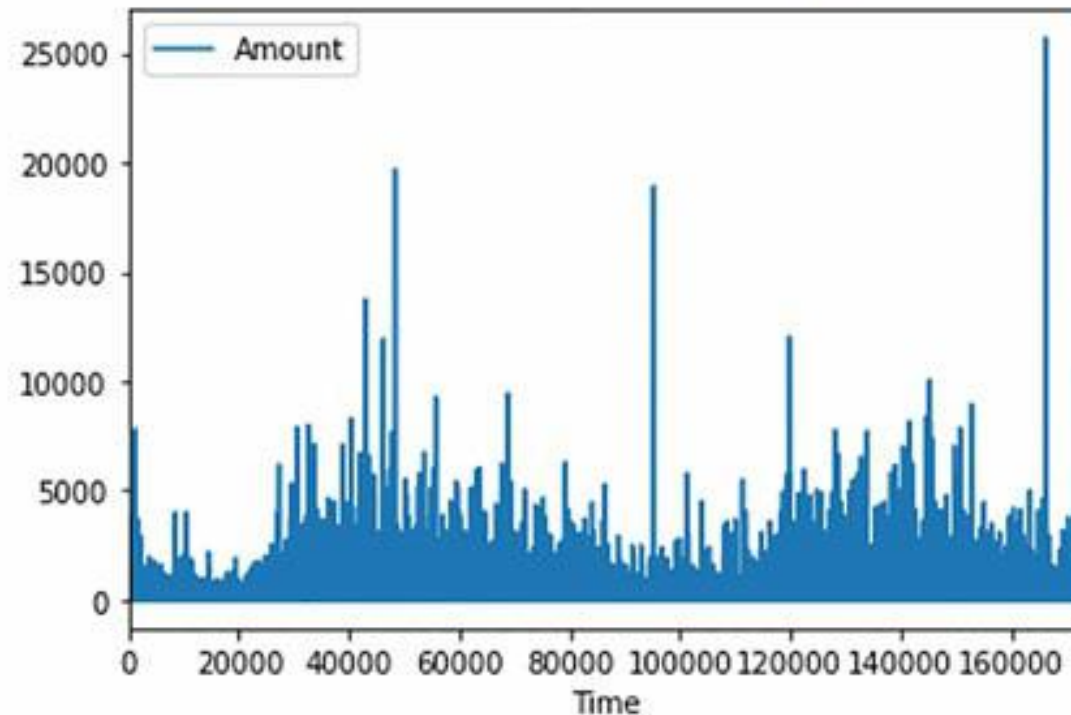
	Time	Amount
0	0.0	149.62
1	0.0	2.69
2	1.0	378.66
3	1.0	123.50
4	2.0	69.99
...
284802	172786.0	0.77
284803	172787.0	24.79
284804	172788.0	67.88
284805	172788.0	10.00
284806	172792.0	217.00

279146 rows × 2 columns

Summed “Amount” Bar Plot


It will make sense if we plot a chart between “Time” vs. “Amount.” We should know that for every “time” what will be the “amounts”

```
df[['Time', 'Amount']].groupby(['Time']).sum().plot()  
<matplotlib.axes._subplots.AxesSubplot at 0x7f7126f5bad0>
```



Class-wise “Amount” vs. “Time” Scatter Plot

Need a plot that
will help us to see
the fraudulent
transactions and
the range for
amounts



Range for Fraudulent Class

The above figure shows the range for the fraudulent transaction to be around [0, 2500]. We can test our hypothesis out and the range of the fraudulent transaction amount

```
f_df.Amount.min(), f_df.Amount.max()
```

```
(0.0, 2125.87)
```


Count of records where “Amount” is zero

How can there be a fraud when the transaction value is zero?

```
df[df["Amount"]==0].shape
```

```
(1825, 31)
```

Count of records where “Amount” is zero

We have a small chunk of records where “Amount” equals to zero. We need to see the breakdown down of “Class” as well and decide whether to keep it or ignore it.

```
df[df["Amount"]==0]["Class"].value_counts()
```

```
0    1798
```

```
1     27
```

```
Name: Class, dtype: int64
```

Count of records where “Amount” is zero

Seeing the breakdown, we can clearly say that it is not significant enough to impact the decision or prediction.

We will see more distributions and later decide to keep it or not.

```
df[df["Amount"]==0]["Class"].value_counts()
```

```
0    1798
```


```
1      27
```

```
Name: Class, dtype: int64
```

Shape of Fraudulent and Non-fraudulent Class

Seeing the above image, we can see the distribution is same or nearly identical if we compare them.


```
fig, ax = plt.subplots(1, 2, figsize=(14,4))  
  
sns.distplot(n_df['Amount'], kde=False, ax=ax[0], color='r')  
ax[0].set_title('Non-Fraudulent: Transaction Amount', fontsize=14)  
  
sns.distplot(f_df['Amount'], kde=False, ax=ax[1], color='b')  
ax[1].set_title('Fraudulent: Transaction Amount', fontsize=14)  
  
plt.show()
```



Shape of Fraudulent and Non-fraudulent Class

KDE will be same as well, which we can confirm by plotting it.


```
fig, ax = plt.subplots(1, 2, figsize=(14,4))  
  
sns.kdeplot(n_df['Amount'], ax=ax[0], color='r')  
ax[0].set_title('Non-Fraudulent: Transaction Amount(KDE)', fontsize=14)  
  
sns.kdeplot(f_df['Amount'], ax=ax[1], color='b')  
ax[1].set_title('Fraudulent: Transaction Amount(KDE)', fontsize=14)  
  
plt.show()
```



“Time” Frequency Distribution for Fraudulent and Non-fraudulent Class

We need to perform something similar for “Time” as well and see the distribution


```
fig, ax = plt.subplots(1, 2, figsize=(14,4))  
  
sns.distplot(n_df['Time'], kde=False, ax=ax[0], color='r')  
ax[0].set_title('Non-Fraudulent: Transaction Time', fontsize=14)  
  
sns.distplot(f_df['Time'], kde=False, ax=ax[1], color='b')  
ax[1].set_title('Fraudulent: Transaction Time', fontsize=14)  
  
plt.show()
```



“Time” Frequency Distribution for Fraudulent and Non-fraudulent Class

Interestingly, the “Time” distribution for Fraudulent transaction is different. We can confirm and analyze after plotting the KDE.


```
fig, ax = plt.subplots(1, 2, figsize=(14,4))  
  
sns.kdeplot(n_df['Time'], ax=ax[0], color='r')  
ax[0].set_title('Non-Fraudulent: Transaction Time(KDE)', fontsize=14)  
  
sns.kdeplot(f_df['Time'], ax=ax[1], color='b')  
ax[1].set_title('Fraudulent: Transaction Time(KDE)', fontsize=14)  
  
plt.show()
```



Interpreting Data and Make Conclusions

Interestingly, the “Time” distribution for Fraudulent transaction is different. We can confirm and analyze after plotting the KDE.

```
fig, ax = plt.subplots(1, 2, figsize=(14,4))  
  
sns.kdeplot(n_df['Time'], ax=ax[0], color='r')  
ax[0].set_title('Non-Fraudulent: Transaction Time(KDE)', fontsize=14)  
  
sns.kdeplot(f_df['Time'], ax=ax[1], color='b')  
ax[1].set_title('Fraudulent: Transaction Time(KDE)', fontsize=14)  
  
plt.show()
```



Interpreting Distributions and Make Conclusions


Non-fraudulent transaction for both “Time” and “Amount” looks the same when compared with the entire population. That is because the size of the Fraudulent dataset is so less the **impact is not significant enough to change the distribution.**

Fraudulent transaction for “Time” looks different, and that has **some significance**, i.e., the density of the distribution is higher at the beginning of time, and there is only one statistical model to the distribution. We need to keep this in mind, when we scale the data and make sure this kind of pattern is not lost.

Non-Fraudulent Summed “Amount” vs. “Time”

As with the previous plots for this class, this is highly similar to the plot for the entire population.


```
n_df[['Time', 'Amount']].groupby(['Time']).sum().plot()  
<matplotlib.axes._subplots.AxesSubplot at 0x7f7131404890>
```



Non-Fraudulent Summed “Amount” vs. “Time”

As we thought, the plot will be different from the hypothesis, and we find the hypothesis to be true.

```
f_df[['Time', 'Amount']].groupby(['Time']).sum().plot()  
<matplotlib.axes._subplots.AxesSubplot at 0x7f7131643250>
```



Total summed amount
for different class

We can see the total amount for each group and removing the zero-amount transaction would not affect the “Amount” distribution, and it won’t change the “Time” distribution as the size of the exception is very low.

```
print("Entire Dataset: " + str(df.Amount.sum()))  
print("Non-Fraudulent Dataset: " + str(n_df.Amount.sum()))  
print("Fraudulent Dataset: " + str(f_df.Amount.sum()))
```

```
Entire Dataset: 25162590.009999998  
Non-Fraudulent Dataset: 25102462.04  
Fraudulent Dataset: 60127.97
```

Removing Zero “Amount” Transaction

We can see the total amount for each group and removing the zero-amount transaction would not affect the “Amount” distribution, and it won’t change the “Time” distribution as the size of the exception is very low.

```
df.drop(df[df.Amount == 0].index, inplace=True)
n_df = df[df["Class"]==0]
f_df = df[df["Class"]==1]
(df.shape, n_df.shape, f_df.shape)

((282982, 31), (282517, 31), (465, 31))
```

1

What is an Imbalanced Dataset?

2

Normalization

3

Scaling

4

Splitting Dataset

5

Handling Imbalance

Scaling, Standard Scaler

Now we need to make all the features same for the comparison.

Standard scalar from the sklearn implements standardization/Z-score normalization.

$$x' = \frac{x_i - \bar{x}}{\sigma}$$

Standard scaling the features

We know from the data description that all the features from V1..., V28 are scaled as those are the output result of a Dimension Reduction Algorithm, i.e., PCA. So, we will only implement standard scaling for “Time” and “Amount.”

```
from sklearn.preprocessing import StandardScaler

std_scaler = StandardScaler()

df['scaled_amount'] = std_scaler.fit_transform(df['Amount'].values.reshape(-1,1))
df['scaled_time'] = std_scaler.fit_transform(df['Time'].values.reshape(-1,1))
```



visualize the change in the distribution

We need to visualize the change in the distribution compared with the original distribution so that we can decide whether to keep or discard it.

```
def compare_kde(coll, col2, name):  
    fig, ax = plt.subplots(2, 2, figsize=(16,7))  
    sns.kdeplot(df[df['Class']==0]['Amount'], ax=ax[0][0], color='r')  
    ax[0][0].set_title('Non-Fraudulent: Transaction Amount(Original)', fontsize=12)  
    sns.kdeplot(df[df['Class']==1]['Time'], ax=ax[0][1], color='b')  
    ax[0][1].set_title('Fraudulent: Transaction Time(Original)', fontsize=12)  
    sns.kdeplot(df[df['Class']==0][coll], ax=ax[1][0], color='r')  
    ax[1][0].set_title('Non-Fraudulent: Transaction Amount('+ name +')', fontsize=12)  
    sns.kdeplot(df[df['Class']==1][col2], ax=ax[1][1], color='b')  
    ax[1][1].set_title('Fraudulent: Transaction Time(RobustScaler)', fontsize=12)  
    plt.show()
```

Correlation heat map after scaling

Interestingly, after scaling, it did not make any visual change, and that is only expected because, “Amount” and “scaled_amount” features correlation equals to one.



Robust Scaling

Robust scaling is similar to standard scaling, but it removes the median and scales the data points according to the IQR.

Robust scaling helps to handle the outliers better than standard scaling. We will also check the distribution and compare it with Standard Scaler.

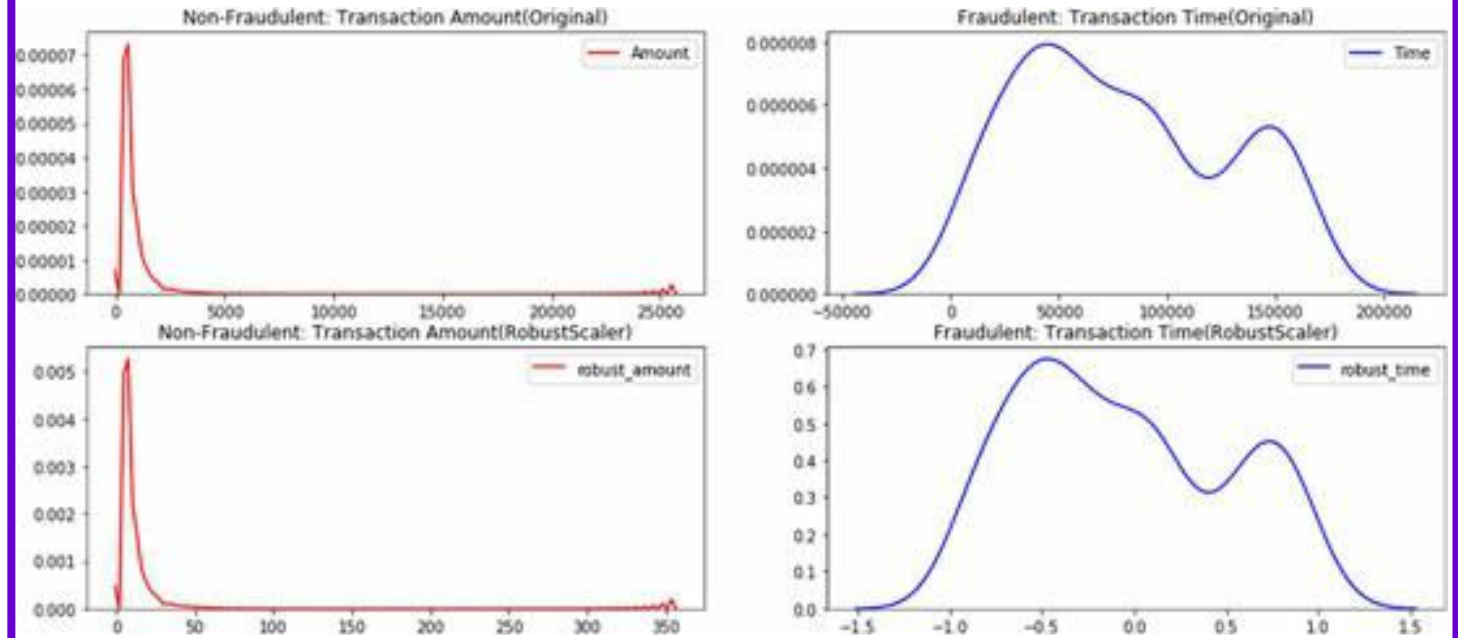
```
from sklearn.preprocessing import RobustScaler

rob_scaler = RobustScaler()

df['robust_amount'] = rob_scaler.fit_transform(df['Amount'].values.reshape(-1,1))
df['robust_time'] = rob_scaler.fit_transform(df['Time'].values.reshape(-1,1))
```

Comparing Robust Scaling

It didn't change that distribution, but the range and the scale are changed. Compared with both the Original and Standard Scaler, it is quite similar



Power Transformer

Power transformer belongs to a family of parametric, monotonic transformations that target to map/project data points from any distribution to as close to a Normal/Gaussian distribution.

This process helps to stabilize the variance of the data and minimize skewness. This method has the highest effects on skewed data.

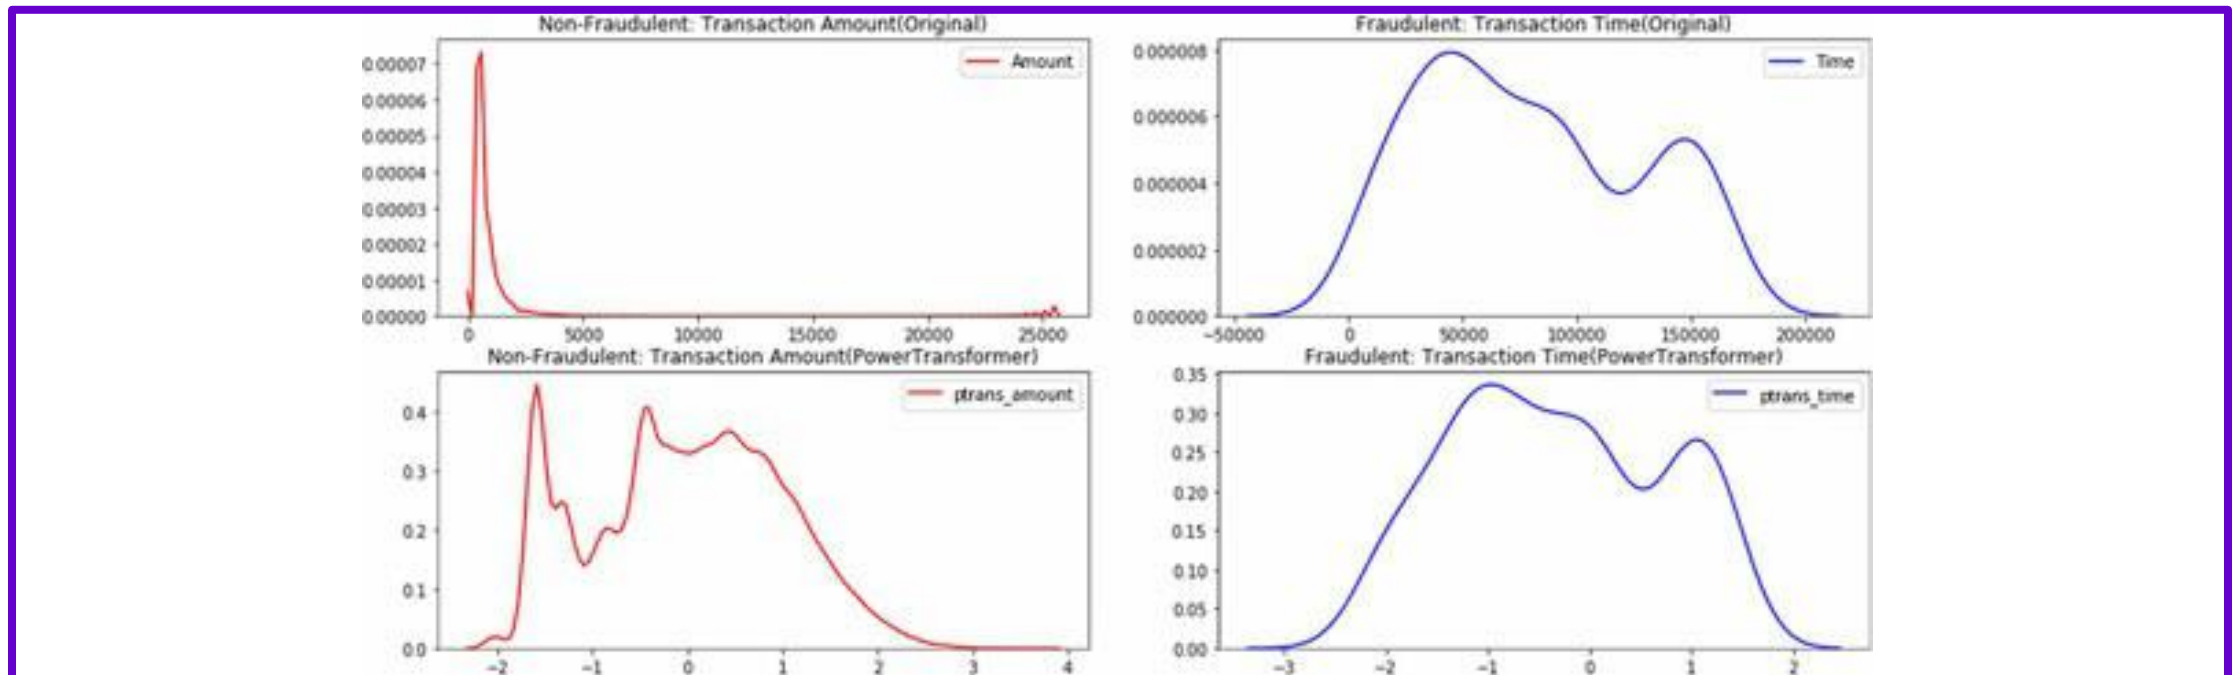
```
from sklearn.preprocessing import PowerTransformer

p_trans = PowerTransformer()

df['ptrans_amount'] = p_trans.fit_transform(df['Amount'].values.reshape(-1,1))
df['ptrans_time'] = p_trans.fit_transform(df['Time'].values.reshape(-1,1))
```


Comparing power transformer

“Amount” feature had quite a lot of change, and it is not similar to a normal distribution but, when compared to “Time” distribution, the change is a lot



Quantile Transformer

Quantile transformer is a non-parametric method to transform the features such that it follows a uniform or a normal distribution. Therefore, for a given feature, this transformation tends to spread out the most frequent values. It also reduces the impact of (marginal) outliers: this is, therefore, a robust pre-processing scheme.

```
from sklearn.preprocessing import QuantileTransformer

q_trans = QuantileTransformer(output_distribution="normal")

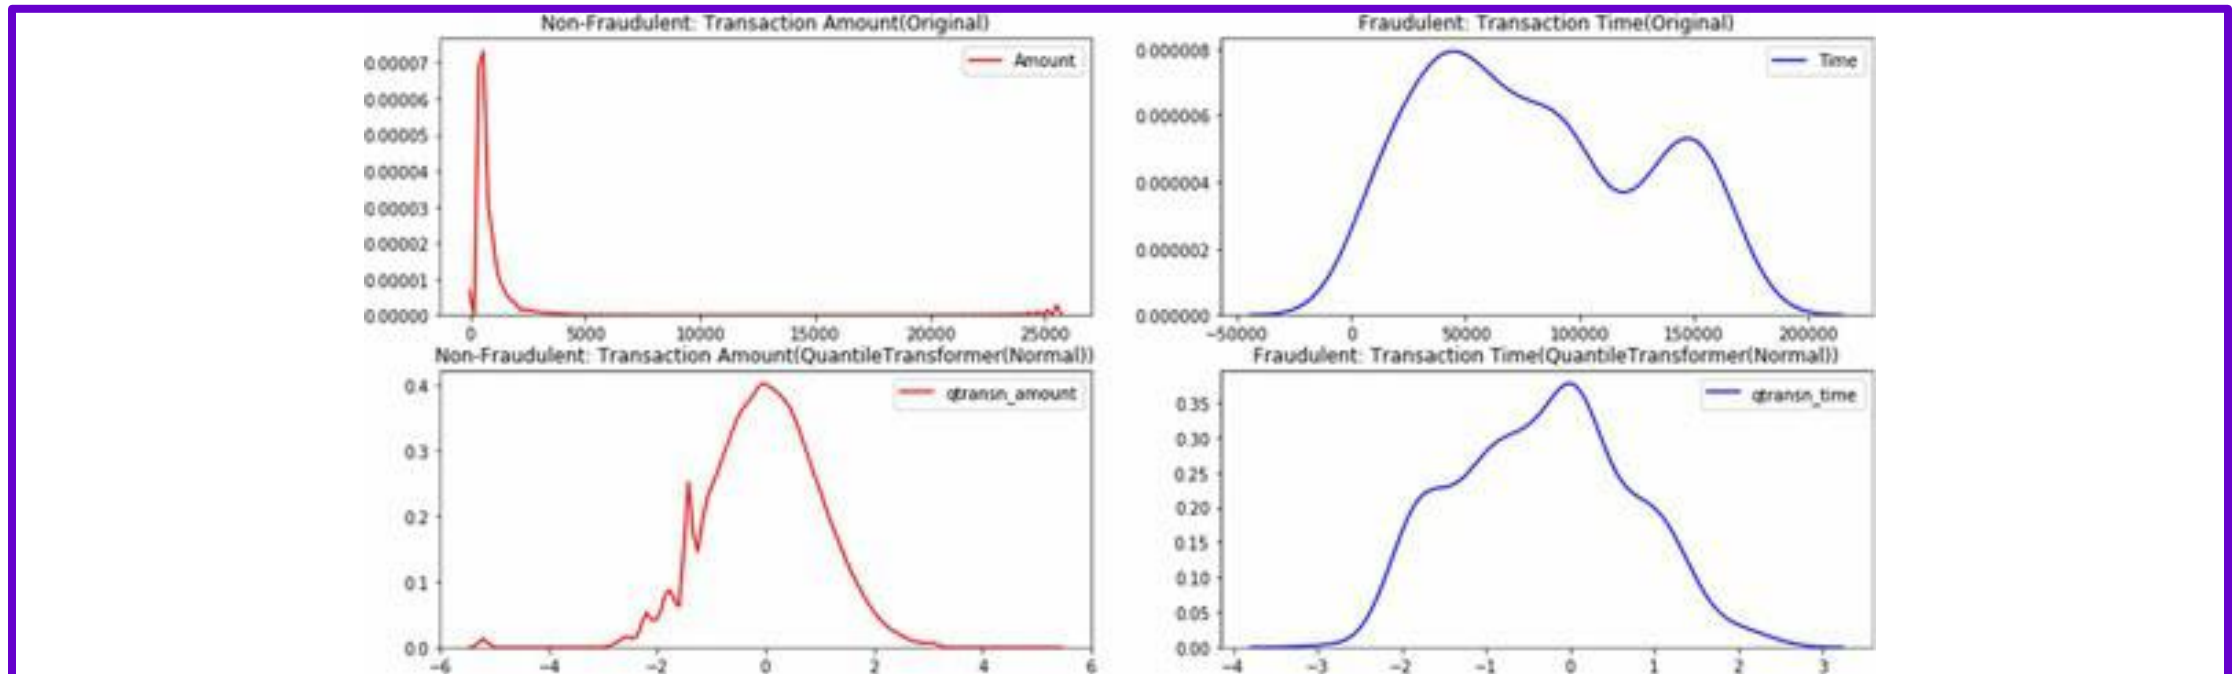
df['qtransn_amount'] = q_trans.fit_transform(df['Amount'].values.reshape(-1,1))
df['qtransn_time'] = q_trans.fit_transform(df['Time'].values.reshape(-1,1))

q_trans = QuantileTransformer(output_distribution="uniform")

df['qtransu_amount'] = q_trans.fit_transform(df['Amount'].values.reshape(-1,1))
df['qtransu_time'] = q_trans.fit_transform(df['Time'].values.reshape(-1,1))
```


Comparing Quantile Transformer -Uniform

We can see that “Amount” is now a uniform distribution, and “Time” is close to a uniform distribution.



Correlation heat map
after applying all
scaling algorithms

Now, we need to
analyze this
correlation matrix
and figure out the
available scaled
feature.



We have one observation that none of the dimensionally reduced features changed its correlation with respect to “Time” and “Amount.”




Code for plotting only scaled features

Analyzing above correlation heat map is very tough as we are not taking V1..., V28 into consideration as of now

```
plt.figure(figsize=(12, 10))
sns.heatmap(df[["ptrans_amount", "qtransn_amount", "qtransu_amount",
               "robust_amount", "scaled_amount", "Amount",
               "Class", "Time", 'scaled_time', "robust_time",
               "qtransu_time", "qtransn_time", "ptrans_time"]].corr(),
            cmap='coolwarm_r', annot=True)
```

Correlation for all scaled features

Analyzing above correlation heat map is very tough as we are not taking V1..., V28 into consideration as of now. So, it is better to plot all the scaled amounts and time separately to analyze them.



Description of scaled “Amount.”

We need to compare “Time,” “Class,” and “Amount” and see an interesting result that has nearly zero change in correlation concerning “Class.”

```
df[["Amount", "scaled_amount", "robust_amount",  
    "qtransu_amount", "qtransn_amount", "ptrans_amount"]].describe().T
```

	count	mean	std	min	25%	50%	75%	max
Amount	282982.0	8.891940e+01	250.824374	0.010000	5.990000	22.490000	78.000000	25691.160000
scaled_amount	282982.0	3.586570e-15	1.000002	-0.354469	-0.330628	-0.264845	-0.043534	102.072560
robust_amount	282982.0	9.225024e-01	3.483188	-0.312179	-0.229135	0.000000	0.770865	356.459797
qtransu_amount	282982.0	5.001364e-01	0.288647	0.000000	0.250250	0.501011	0.748822	1.000000
qtransn_amount	282982.0	-4.655992e-03	1.019305	-5.199338	-0.673702	0.001756	0.671178	5.199338
ptrans_amount	282982.0	2.142510e-14	1.000002	-2.050775	-0.733013	0.030784	0.750683	3.649005

Description of scaled “Time.”

“Time” also reflects the same characteristics as “Amount,” but the change is drastic for some algorithm for the future “Amount.”

```
df[["Time", 'scaled_time', "robust_time",  
    "qtransu_time", "qtransn_time", "ptrans_time"]].describe().T
```

	count	mean	std	min	25%	50%	75%	max
Time	282982.0	9.484896e+04	47482.459589	0.000000	54251.250000	84707.500000	139363.750000	172792.000000
scaled_time	282982.0	6.155099e-15	1.000002	-1.997561	-0.855006	-0.213584	0.937501	1.641515
robust_time	282982.0	1.191536e-01	0.557879	-0.995242	-0.357835	0.000000	0.642165	1.034918
qtransu_time	282982.0	4.995409e-01	0.288562	0.000000	0.248568	0.499854	0.749264	1.000000
qtransn_time	282982.0	3.381409e-03	1.000297	-5.199338	-0.671536	0.005473	0.679379	5.199338
ptrans_time	282982.0	1.981191e-14	1.000002	-2.436283	-0.809483	-0.143207	0.928790	1.534947

Being an ML practitioner

Like we have seen, **there is not much information** we found out from this scaling section

Being an ML practitioner, **we need to do the same thing again and again**, and every work depends on the **experience**.

but, when we will go ahead with the next section, i.e., handling the imbalanced data, **we will see a huge improvement** with the results from scaling too

1

What is an Imbalanced Dataset?

2

Normalization

3

Scaling

4

Splitting Dataset

5

Handling Imbalance

Splitting Dataset

We need to split the data for many reasons. One of the major reasons is for testing the model

```
from sklearn.model_selection import train_test_split
import numpy as np

print('No Frauds', round(df['Class'].value_counts()[0]/len(df) * 100,2), '% of the dataset')
print('Frauds', round(df['Class'].value_counts()[1]/len(df) * 100,2), '% of the dataset')

X = df.drop('Class', axis=1)
y = df['Class']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=2)

# See if both the train and test label distribution are similarly distributed
train_unique_label, train_counts_label = np.unique(y_train, return_counts=True)
test_unique_label, test_counts_label = np.unique(y_test, return_counts=True)

print('\nLabel Distributions: \n')
print(train_counts_label/ len(original_ytrain))
print(test_counts_label/ len(original_ytest))
print("\nTrain:")
print('No Frauds', round(len(y_train[y_train==0])/len(X_train) * 100,2), '% of the dataset')
print('Frauds', round(len(y_train[y_train==1])/len(X_train) * 100,2), '% of the dataset')
print("\nTest:")
print('No Frauds', round(len(y_test[y_test==0])/len(X_test) * 100,2), '% of the dataset')
print('Frauds', round(len(y_test[y_test==1])/len(X_test) * 100,2), '% of the dataset')
```

Split Data Statistics (Traditional)

We need to make the fraud percentage the same or similar to the original fraud percentage. As we are dealing with an extremely small percentage of fraud data, our target should be the same.

No Frauds 99.84 % of the dataset
Frauds 0.16 % of the dataset

Label Distributions:

```
[0.99829936 0.00169622]  
[0.99858647 0.0014312 ]
```

Train:

No Frauds 99.83 % of the dataset
Frauds 0.17 % of the dataset

Test:

No Frauds 99.86 % of the dataset
Frauds 0.14 % of the dataset

Split Data Statistics (StratifiedKFold)

There is a solution to this problem that is stratified K-fold, it is a similar concept like a cross validation K-fold technique. It gives us the same distribution as the original distribution

```
from sklearn.model_selection import StratifiedKFold
import numpy as np

print('No Frauds', round(df['Class'].value_counts()[0]/len(df) * 100,2), '% of the dataset')
print('Frauds', round(df['Class'].value_counts()[1]/len(df) * 100,2), '% of the dataset')

X = df.drop('Class', axis=1)
y = df['Class']

skf = StratifiedKFold(n_splits=10, random_state=None, shuffle=False)

for train_index, test_index in skf.split(X, y):
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

# See if both the train and test label distribution are similarly distributed
train_unique_label, train_counts_label = np.unique(y_train, return_counts=True)
test_unique_label, test_counts_label = np.unique(y_test, return_counts=True)

print('\nLabel Distributions: \n')
print(train_counts_label/ len(original_ytrain))
print(test_counts_label/ len(original_ytest))
print("\nTrain:")
print('No Frauds', round(len(y_train[y_train==0])/len(X_train) * 100,2), '% of the dataset')
print('Frauds', round(len(y_train[y_train==1])/len(X_train) * 100,2), '% of the dataset')
print("\nTest:")
print('No Frauds', round(len(y_test[y_test==0])/len(X_test) * 100,2), '% of the dataset')
print('Frauds', round(len(y_test[y_test==1])/len(X_test) * 100,2), '% of the dataset')
```

Splitting Dataset (StratifiedKFold)

From the above code snippet, we are expecting the training and testing dataset with the class distribution like the original one

```
No Frauds 99.84 % of the dataset  
Frauds 0.16 % of the dataset
```

```
Label Distributions:
```

```
[1.12315249 0.00185082]  
[0.49916955 0.00081278]
```

```
Train:
```

```
No Frauds 99.84 % of the dataset  
Frauds 0.16 % of the dataset
```

```
Test:
```

```
No Frauds 99.84 % of the dataset  
Frauds 0.16 % of the dataset
```

Splitting Dataset

We got the expected output, i.e., the percentage of the class distribution is the same for the train, test, and the original dataset. Now, we can go ahead with handling imbalance characteristic of the dataset.

1

What is an Imbalanced Dataset?

2

Normalization

3

Scaling

4

Splitting Dataset

5

Handling Imbalance

Handling Imbalance: Oversampling

In oversampling, we will generally create synthetic data points for the class, which has low counts.

So, for oversampling, we have to use the test data from the original data, because after performing oversampling, if we split the data that won't be correct, and we will get a wrong result.

Handling Imbalance: Under-sampling

for under-sampling, we can use the original dataframe to make a sub-sample out of it.

But, in this case, we will only use the train test dataset.

What is a sub-sample?

Sub-sample just means a part of the original dataset, but in this scenario, the subsample will be 50-50.

We will take the entire Fraudulent class and under-sample the non-fraudulent class so that the ratio is 50-50.

Advantages of sub-subsample

Overfitting is a common problem when there is an imbalance because it assumes that, in most cases, there are non-fraudulent transactions. But for a subsample, this problem will be less prone to overfitting.

Correlation with the class will improve drastically as our subsample will have no imbalance problem. Although we don't know what the "V" features stand for (but, my assumption is "Vector"), it will be useful to understand how each of these features influences the result with the class (Fraud or No Fraud) by having an imbalance dataframe we are not able to see the true correlations between the class and features


What is Resampling

Overfitting is a common problem when there is an imbalance because **it assumes that, in most cases, there are non-fraudulent transactions**. But for a subsample, this problem will be less prone to overfitting.

Correlation with the class will improve drastically as our subsample will have no imbalance problem. Although we don't know what the "V" features stand for (but, my assumption is "Vector"), it will be useful to understand how each of these features influences the result with the class (Fraud or No Fraud) by having an imbalance dataframe we are not able to see the true correlations between the class and features

What is Resampling

Resampling is a technique to handle imbalance data by creating a sub-sampled dataset from the original data.



Resampling Train Datas

We need to make sure that we are **using only the train dataset** and not using the test dataset at all. The test dataset will only be used to validate the model and see how good it is.

```
train_df = X_train.copy()  
train_df[ 'Class' ] = y_train  
train_df.shape
```

```
(254685, 41)
```

Train Dataframe Class Distribution

Within that training dataframe, there will be a mixture of fraud and non-fraud datasets in the ratio of the Original dataset as we have used the Stratified K-fold technique

```
train_df[ 'Class' ].value_counts()
```

```
0      254266
```

```
1       419
```

```
Name: Class, dtype: int64
```


PCA plot for Original dataframe

Before going ahead, we need to find a way to visualize the 41 column features. We can use a dimensionality reduction algorithm like Principal Component Analysis to reduce to 2 principal components and plot in a 2D plot

```
from sklearn.decomposition import PCA

pca = PCA(n_components=2)
X = pca.fit_transform(X_train)

plot_2d_space(X, y_train, 'Imbalanced dataset (2 PCA components)')
```



Function to plot 2 component PCA

As we will plot a 2D graph, again and again, we are using a function for that. Then a simple function call is sufficient to plot it

```
def plot_2d_space(X, y, label='Class'):  
    colors = ['#1F77B4', '#FF7F0E']  
    markers = ['o', 's']  
    for l, c, m in zip(np.unique(y), colors, markers):  
        plt.scatter(  
            X[y==l, 0],  
            X[y==l, 1],  
            c=c, label=l, marker=m  
        )  
    plt.title(label)  
    plt.legend(loc='upper right')  
    plt.show()
```

Random Undersampling

Random under-sampling will remove the majority class data points such that it is equivalent/equal/proportional to the minority class

```
# Class count
count_class_0, count_class_1 = train_df.Class.value_counts()

# Divide by class
train_df_0 = train_df[train_df['Class'] == 0]
train_df_1 = train_df[train_df['Class'] == 1]
```

Random undersampling

Now, this dataset is equidistributed, and it is likely to give a better correlation and model results, as we have hypothesized before.

can arise multiple problems due to information loss

```
train_df_0_under = train_df_0.sample(count_class_1)
train_df_under = pd.concat([train_df_0_under, train_df_1], axis=0)

print('Random under-sampling:')
print(train_df_under.Class.value_counts())

sns.countplot('Class', data=train_df_under)
```


Random under-sampling:

1 419

0 419

Name: Class, dtype: int64


<matplotlib.axes._subplots.AxesSubplot at 0x1a31c29fd0>



PCA plot for random undersampling


we can see the homogeneity in the data points of different classes. With this new homogeneous dataset, we are expecting a good correlation between all the features and primarily with the feature “class.”

```
from sklearn.decomposition import PCA  
  
pca = PCA(n_components=2)  
X = pca.fit_transform(train_df_under.drop(columns=["Class"]))  
  
plot_2d_space(X, train_df_under["Class"], 'Undersampled dataset (2 PCA components)')
```




Correlation for Random Under Sampled data

We can see a drastic change with the correlation matrix for mostly all the features. Some features are now positively and negatively correlated



Correlation for
Random Under
Sampled data...

There is one more
interesting
observation that after
random under-
sampling, all the
scaled features and
their respective
original features have
the same or similar
correlation



Random over-sampling

Similar to random under-sampling, we have random over-sampling where we will oversample the minor class, i.e., in this case, it will be the fraudulent transaction dataset.

```
train_df_1_over = train_df_1.sample(count_class_0, replace=True)
train_df_over = pd.concat([train_df_0, train_df_1_over], axis=0)

print('Random over-sampling:')
print(train_df_over.Class.value_counts())

sns.countplot('Class', data=train_df_over)
```


Random over-sampling:

1 254266

0 254266

Name: Class, dtype: int64

<matplotlib.axes._subplots.AxesSubplot at 0x1a336eee90>




PCA plot for random over-sampling

The Oversampled data set shows the same plot as the original dataset plot.

So that it means the correlation heat map will be similar to the original one?? No


```
from sklearn.decomposition import PCA  
  
pca = PCA(n_components=2)  
X = pca.fit_transform(train_df_over.drop(columns=["Class"]))  
  
plot_2d_space(X, train_df_over["Class"], 'Oversampled dataset (2 PCA components)')
```



PCA plot for random over-sampling


So that it means the correlation heat map will be similar to the original one??

But unfortunately, no, it will be extremely different because this oversampled dataset is equidistributed so it will have a better correlation and will be similar to the randomly under-sampled data set.




Tomek's Links for Undersampling: Step 1

The data points are plotted in n-dimensions(n-features) in this case for simplicity, we are using 2-axis and differentiate them with respect to class (here, there are 2 classes only).




Tomek's Links for Undersampling: Step 2

We find the nearest neighbor where $k=1$ where the minor and the major classes are coupled.




Tomek's Links for Undersampling: Step 3

After we find the
coupled data, we
will remove the
major class for
under-sampling.



Tomek's Links for Undersampling

The above illustrations give us a simple working of Tomek's Link. This helps us to define a good decision boundary and remove some of the data points from the majority class.



Tomek's Links for Undersampling

We can clearly state that there won't be any change from the original data as the change is not significant enough. Let's plot the principal components and see how it looks.

```
from imblearn.under_sampling import TomekLinks

tkl = TomekLinks(sampling_strategy='auto', n_jobs=-1)
X_tkl, y_tkl = tkl.fit_sample(X_train, y_train)

train_df_tkl = X_tkl
train_df_tkl['Class'] = y_tkl

print('Tomek links under-sampling:')
print(train_df_tkl.Class.value_counts())
sns.countplot('Class', data=train_df_tkl)
```


Tomek links under-sampling:

0 254204

1 418


Name: Class, dtype: int64

<matplotlib.axes._subplots.AxesSubplot at 0x1a5e531890>




PCA plot for Tomek's Link

We can see it is the same distribution as the original dataframe. We can confirm that the heat map did not change at all, so this technique cannot be used for under-sampling.




Synthetic Minority Oversampling Technique: Step 1

Let us plot the data
points in
dimensional
space (n-features)
and classify them as
per the "Class."




Synthetic Minority Oversampling Technique: Step 2

Create a pattern/boundary of the minority dataset with the data points. We can create a space using the boundary of the minority data and add synthetic values within the boundary where the nearest neighbor is also of the same class.



Synthetic Minority Oversampling Technique: Step 3

Plot those points
and which meets
step 2 conditions
and assign them
the minor class.



Synthetic Minority Oversampling Technique

We can see that it
equidistributed,
and we have the
same number of
records for both the
classes.

```
from imblearn.over_sampling import SMOTE

sm = SMOTE()
X_sm, y_sm = sm.fit_sample(X_train, y_train)

train_df_sm = X_sm
train_df_sm['Class'] = y_sm

print('SMOTE over-sampling:')
print(train_df_sm.Class.value_counts())
sns.countplot('Class', data=train_df_sm)
```


SMOTE over-sampling:

1 254266

0 254266

Name: Class, dtype: int64


<matplotlib.axes._subplots.AxesSubplot at 0x1a6b2f3cd0>



PCA plot for SMOTE


We can see there is a difference in the distribution. But it is not significant enough(it seems so)

```
from sklearn.decomposition import PCA  
  
pca = PCA(n_components=2)  
X = pca.fit_transform(train_df_sm.drop(columns=["Class"]))  
  
plot_2d_space(X, train_df_sm["Class"], 'SMOTE over-sampling (2 PCA components)')
```



Correlation Heatmap for SMOTE dataset

Interestingly, we
can see a similar
heat map for both
Random
Oversampling and
SMOTE



Function for plot Box-and-whiskers plot for Positive Correlated Data


From the last section, we have seen the features V2, V4, V11, and V19 are positively correlated.

We will write a function that will help us to plot all the boxplots to compare them.

```
def plot_pos_box(df):  
    f, axes = plt.subplots(ncols=4, figsize=(20,7))  
    colors = ["#0101DF", "#DF0101"]  
  
    sns.boxplot(x="Class", y="V11", data=df, palette=colors, ax=axes[0])  
    axes[0].set_title('V11 vs Class Positive Correlation')  
  
    sns.boxplot(x="Class", y="V4", data=df, palette=colors, ax=axes[1])  
    axes[1].set_title('V4 vs Class Positive Correlation')  
  
    sns.boxplot(x="Class", y="V2", data=df, palette=colors, ax=axes[2])  
    axes[2].set_title('V2 vs Class Positive Correlation')  
  
    sns.boxplot(x="Class", y="V19", data=df, palette=colors, ax=axes[3])  
    axes[3].set_title('V19 vs Class Positive Correlation')  
  
    plt.show()
```


Original training data

We can see that the original training dataframe has an extremely large number of outliers for class 0, and that is very bad for modeling, so we went for resampling the dataset




Random under-sampling

This plot looks way better as there are a smaller number of outliers, and this can be used for training machine learning models.




Random over-sampling

As this shows so many outliers, we cannot use this dataset for modeling. So, we will be ignoring this dataset.




Tomek's Link under-sampling

This plot is exactly similar to the original training dataset, as there were only 62 records that were removed from the majority class.



SMOTE

SMOTE is also an oversampling technique so that it will look similar to Random Oversampling.



Function for plot Box-and-whiskers plot for Negative Correlated Data

Similarly, we will write a function for negatively correlated values to plot box-and-whiskers and find the optimal resampling technique

```
def plot_neg_box(df):  
    f, axes = plt.subplots(ncols=5, figsize=(28,7))  
    colors = ["#0101DF", "#DF0101"]  
  
    sns.boxplot(x="Class", y="V17", data=df, palette=colors, ax=axes[0])  
    axes[0].set_title('V17 vs Class Negative Correlation')  
  
    sns.boxplot(x="Class", y="V14", data=df, palette=colors, ax=axes[1])  
    axes[1].set_title('V14 vs Class Negative Correlation')  
  
    sns.boxplot(x="Class", y="V12", data=df, palette=colors, ax=axes[2])  
    axes[2].set_title('V12 vs Class Negative Correlation')  
  
    sns.boxplot(x="Class", y="V16", data=df, palette=colors, ax=axes[3])  
    axes[3].set_title('V16 vs Class Negative Correlation')  
  
    sns.boxplot(x="Class", y="V10", data=df, palette=colors, ax=axes[4])  
    axes[4].set_title('V10 vs Class Negative Correlation')  
  
    plt.show()
```


Function for plot Box-and-whiskers plot for Scaled Amount

Now we have to select which scaling technique to use for the modeling purpose. We have selected only two resampling techniques from the previous section. With that dataframe, we will be looking at the scaled data and select the optimal algorithm for the scaled feature ("Time" and "Amount").

```
def plot_scaled_amt(df):  
    f, axes = plt.subplots(ncols=5, figsize=(28,7))  
    colors = ["#0101DF", "#DF0101"]  
  
    sns.boxplot(x="Class", y="scaled_amount", data=df, palette=colors, ax=axes[0])  
    axes[0].set_title('Standard Scaling (Amount)')  
  
    sns.boxplot(x="Class", y="robust_amount", data=df, palette=colors, ax=axes[1])  
    axes[1].set_title('Robust Scaling (Amount)')  
  
    sns.boxplot(x="Class", y="ptrans_amount", data=df, palette=colors, ax=axes[2])  
    axes[2].set_title('Power Transformer (Amount)')  
  
    sns.boxplot(x="Class", y="qtransn_amount", data=df, palette=colors, ax=axes[3])  
    axes[3].set_title('Quartile Transformer -Normal (Amount)')  
  
    sns.boxplot(x="Class", y="qtransu_amount", data=df, palette=colors, ax=axes[4])  
    axes[4].set_title('Quartile Tranformer -Uniform (Amount)')  
  
    plt.show()
```

Function for plot Box-and-whiskers plot for Scaled Time


Now we have to select which scaling technique to use for the modeling purpose. We have selected only two resampling techniques from the previous section. With that dataframe, we will be looking at the scaled data and select the optimal algorithm for the scaled feature ("Time" and "Amount").

```
def plot_scaled_time(df):  
    f, axes = plt.subplots(ncols=5, figsize=(28,7))  
    colors = ["#0101DF", "#DF0101"]  
  
    sns.boxplot(x="Class", y="scaled_time", data=df, palette=colors, ax=axes[0])  
    axes[0].set_title('Standard Scaling (Time)')  
  
    sns.boxplot(x="Class", y="robust_time", data=df, palette=colors, ax=axes[1])  
    axes[1].set_title('Robust Scaling (Time)')  
  
    sns.boxplot(x="Class", y="ptrans_time", data=df, palette=colors, ax=axes[2])  
    axes[2].set_title('Power Transformer (Time)')  
  
    sns.boxplot(x="Class", y="qtransn_time", data=df, palette=colors, ax=axes[3])  
    axes[3].set_title('Quartile Transformer -Normal (Time)')  
  
    sns.boxplot(x="Class", y="qtransu_time", data=df, palette=colors, ax=axes[4])  
    axes[4].set_title('Quartile Tranformer -Uniform (Time)')  
  
    plt.show()
```

Time:

Original Training Data for Box plot We


We can see for “Quantile Transformer -Normal” there is skewness and outliers in the plot. We now have to compare this entire plot with Random Undersampling and SMOTE



Time:

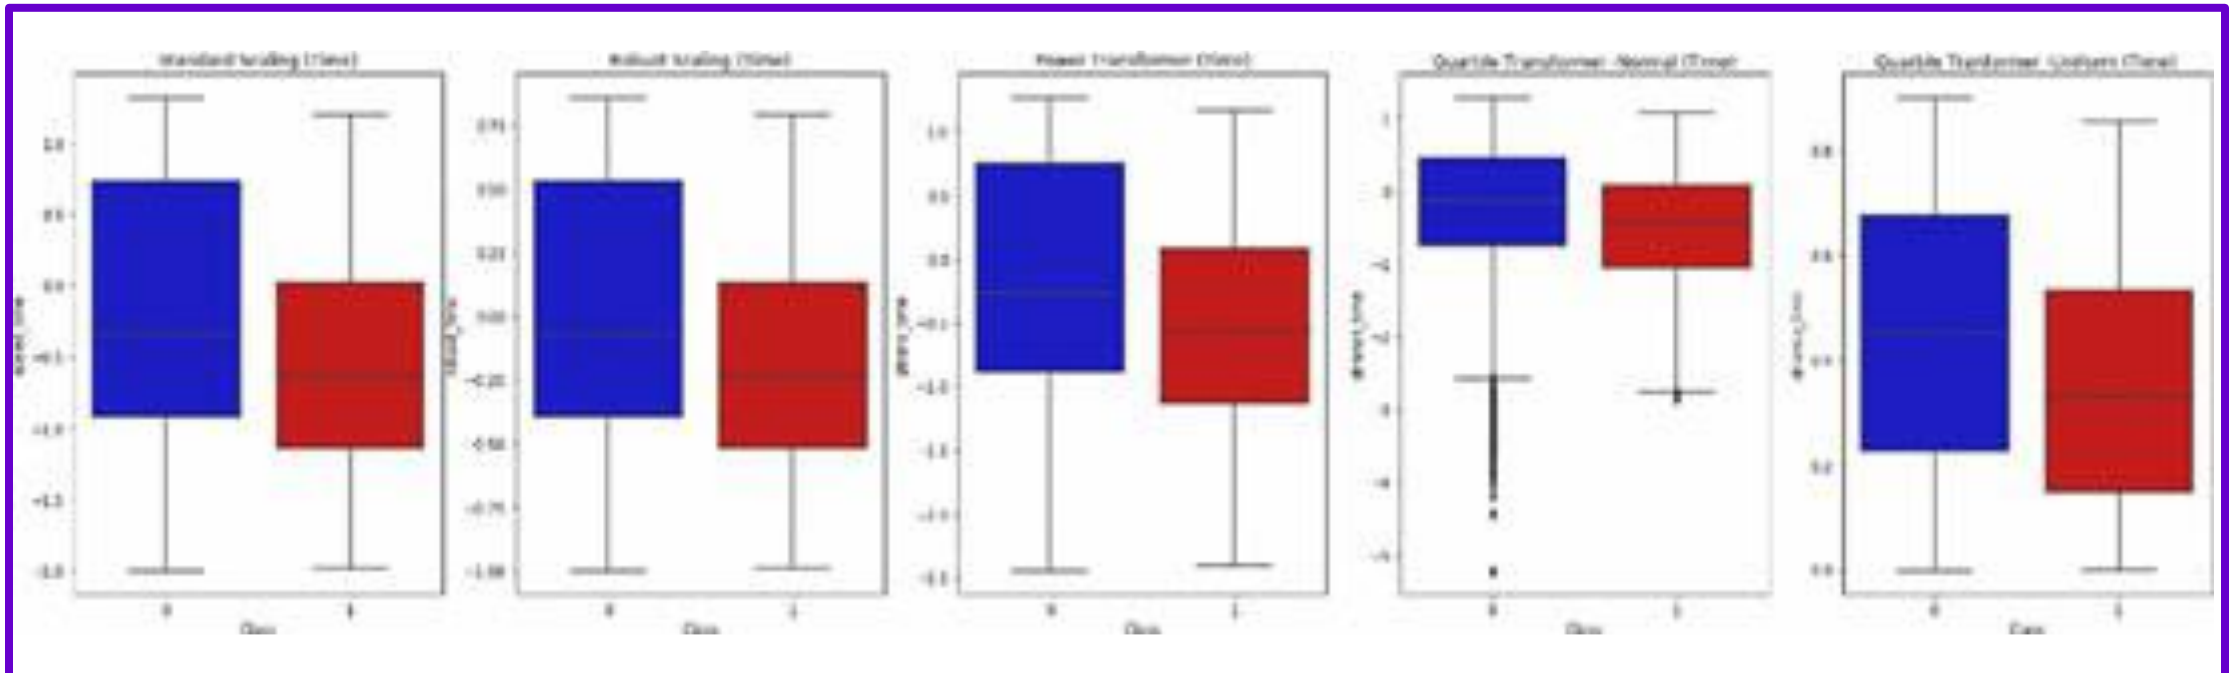
Original Training Data for Box plot We

For this dataset, we can see the same thing that for “Quantile Transformer - Normal,” there are outliers in the plot. Apart from that, all looks good for use. Previously, we have seen that “Quantile Transformer -Uniform” actually changed the distribution so that can be ignored as well.



Time: SMOTE for Box plot


For the SMOTE dataframe, we see the same result, so, for the feature “Time,” we can choose any algorithm apart from “Quantile Transformer.”



Amount:

Original Training Data for Box plot We

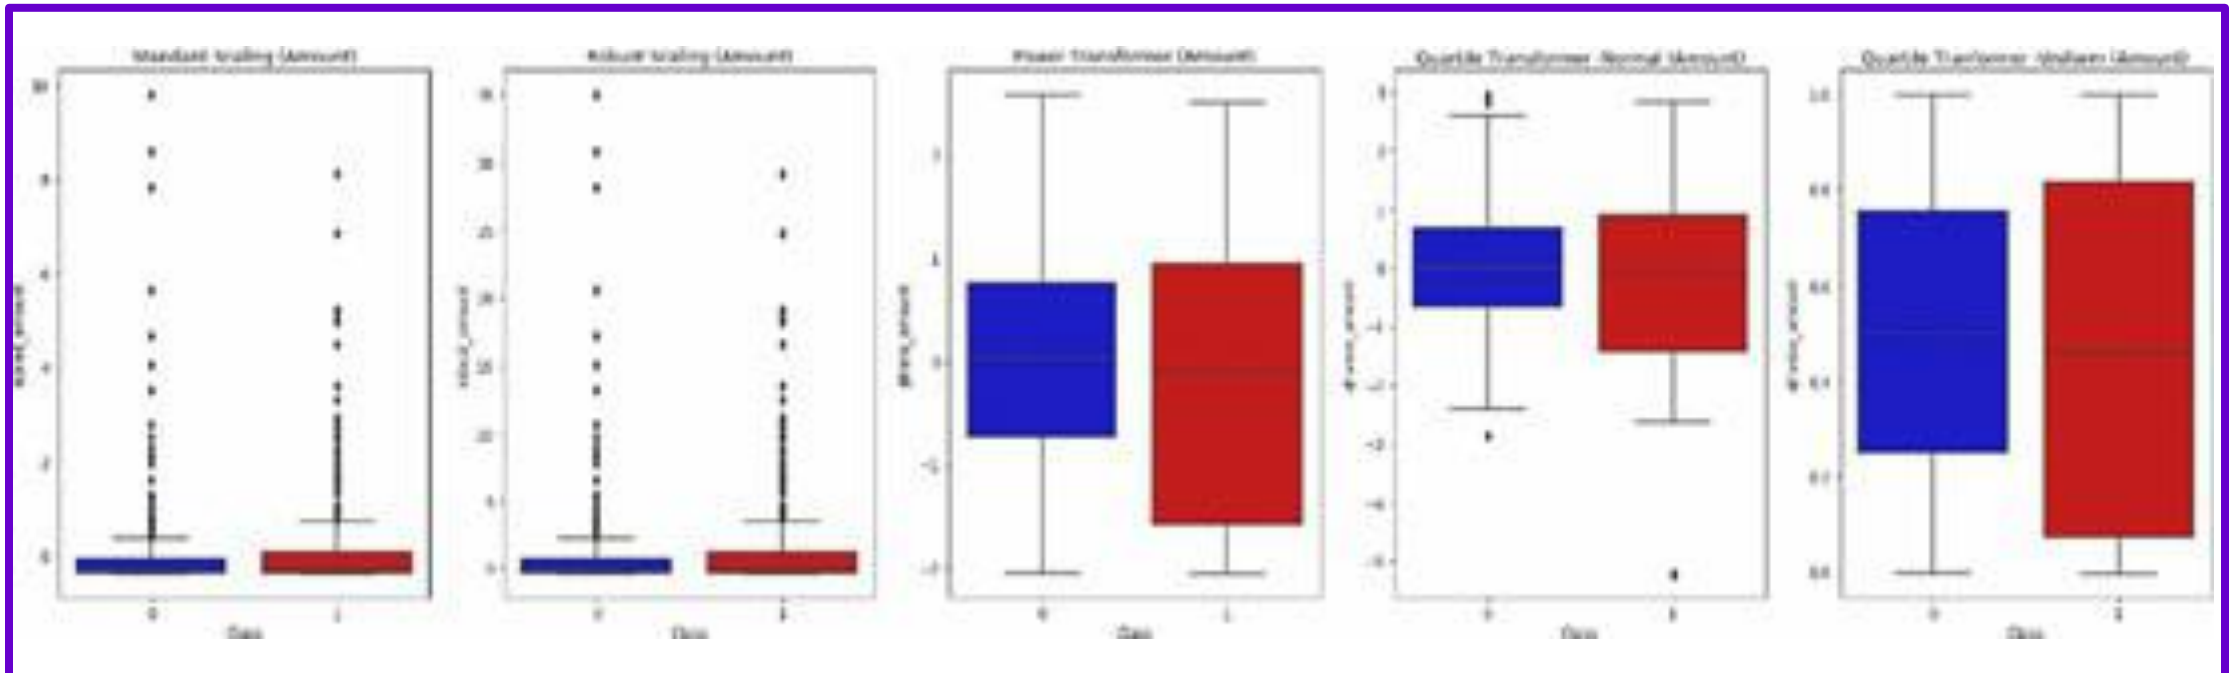
Here, for the amount, we cannot ignore or eliminate the outliers as they will give us a lot of information. We will try to stick with the original distribution with some minor changes.



Amount:

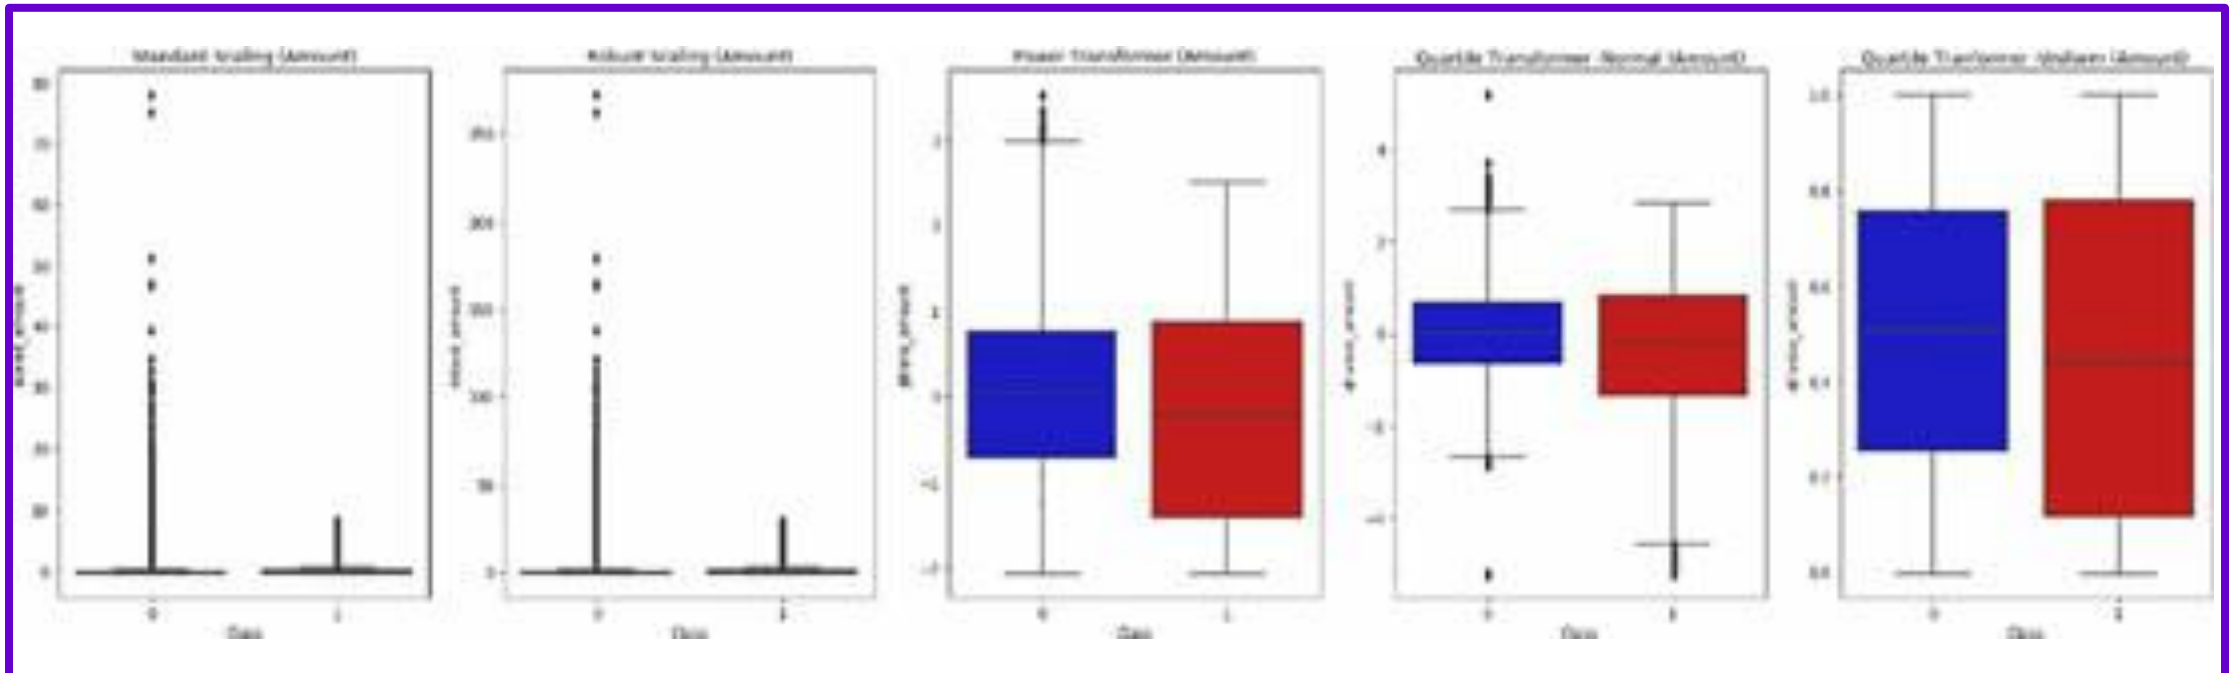
Random under sampling for Box plot

We can see outliers have reduced, and Power Transformer looks good in terms of distribution. But we will also look at Robust Scaler as it more or less gives us the original distribution. “Quantile Transformer - Normal” looks good, but it has already performed badly for the future “Time,” so we will ignore it



Amount: SMOTE for Box plot

SMOTE plot shows a similar thing like the original training dataset; hence we will stick with our last decision.



Find the best(Optimal) data for the model

SMOTE Dataset 1 - Robust Scaled Time, Robust Scaled Amount, V2, V4, V11, and V19 (Positive), V17, V14, V12, V16 and V10 (Negative)

SMOTE Dataset 2 - Power Transformer Time, Power Transformer Amount, V2, V4, V11, and V19 (Positive), V17, V14, V12, V16 and V10 (Negative)

Random Under-sample Dataset 1 - Robust Scaled Time, Robust Scaled Amount, V2, V4, V11, and V19 (Positive), V17, V14, V12, V16 and V10 (Negative)

Random Under-sample Dataset 2 - Power Transformer Time, Power Transformer Amount, V2, V4, V11, and V19 (Positive), V17, V14, V12, V16 and V10 (Negative)

6

Logistic Regression

7

Decision Tree

8

Support Vector Machine (SVM)

9

Data Preparation & Metric Trap

10

Training & Evaluation

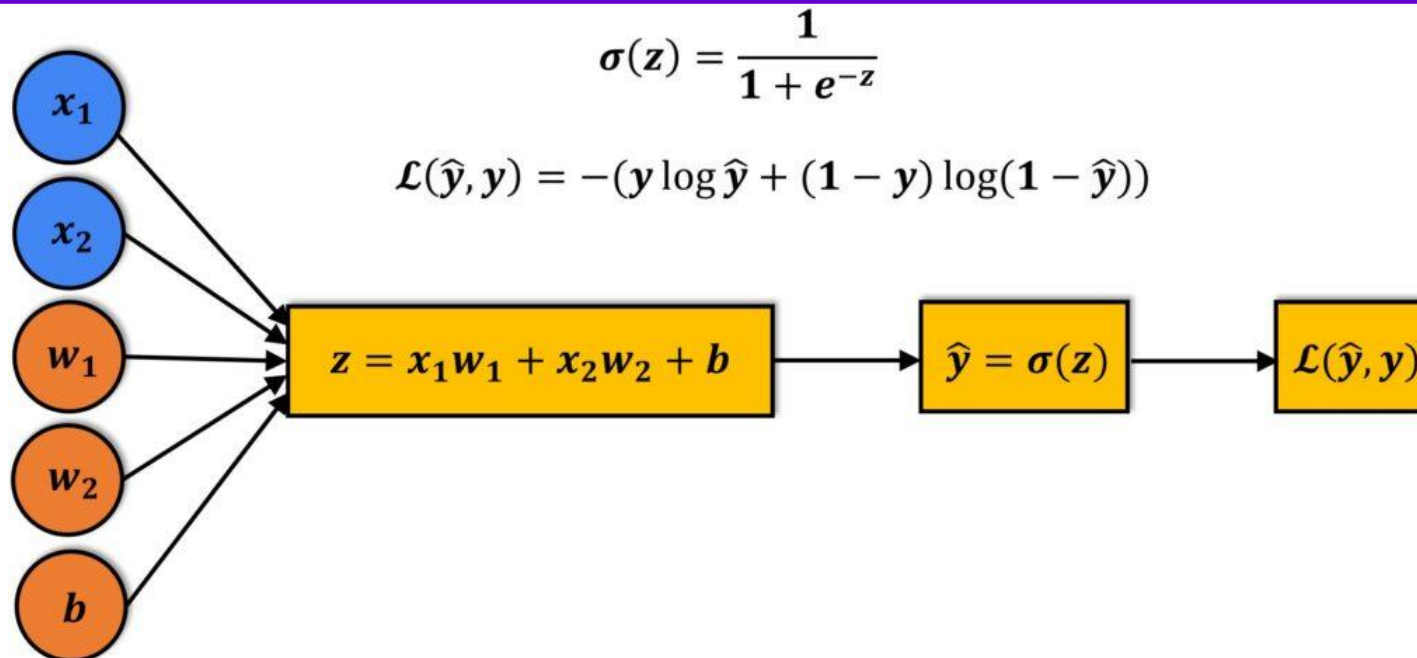
Logistic Regression Model

Logistic Regression is a predictive analysis which is used to explain the data and **relationship between** one dependent **binary variable** and **one or more nominal, ordinal, interval or ratio-level independent variables**.

$$p(x) = \frac{e^{(\beta_0 + \beta_1 x_1 + \dots + \beta_k x_k)}}{1 + e^{(\beta_0 + \beta_1 x_1 + \dots + \beta_k x_k)}}$$


Logistic Regression Model

A linear equation (z) is given to a sigmoidal activation function (σ) to predict the output (\hat{y}).



Logistic Regression Vs Linear Regression

Unlike linear regression, we get an 'S' shaped curve in logistic regression.



6

Logistic Regression

7

Decision Tree

8

Support Vector Machine (SVM)

9

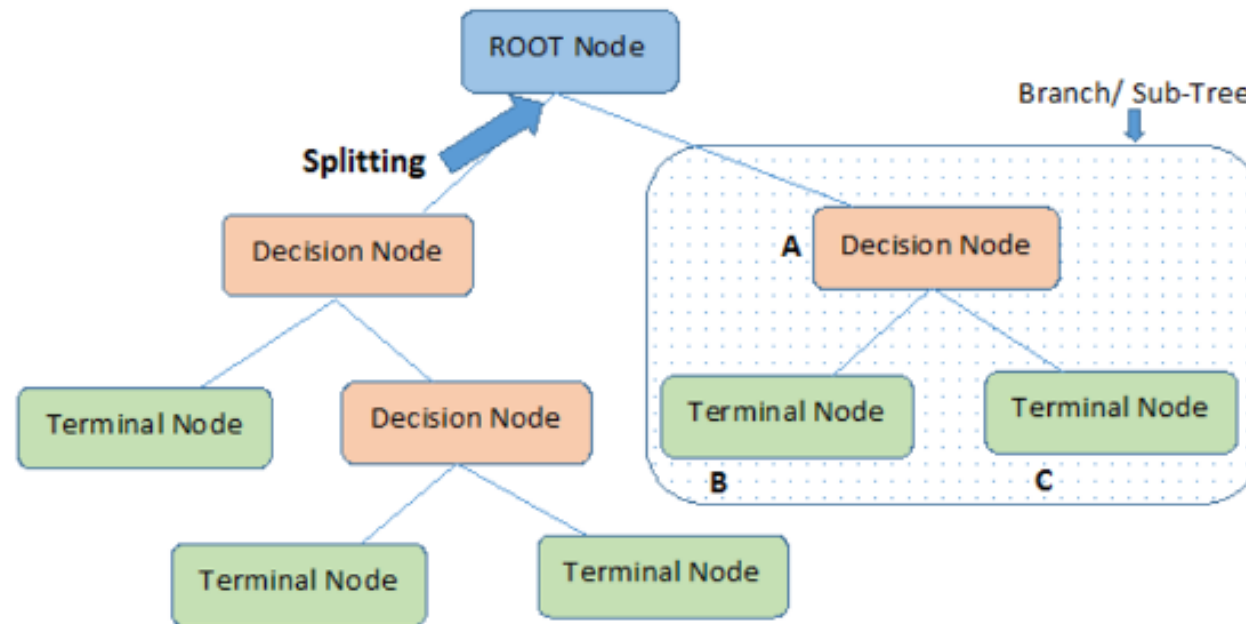
Data Preparation & Metric Trap

10

Training & Evaluation

Decision Tree Model

A decision tree is a predictive model that uses a flowchart-like structure to make decisions based on input data. It divides data into branches and assigns outcomes to leaf nodes.




Example of Decision Tree

Decision trees are upside down which means the root is at the top and then this root is split into various several nodes.

Day	Weather	Temperature	Humidity	Wind	Play?
1	Sunny	Hot	High	Weak	No
2	Cloudy	Hot	High	Weak	Yes
3	Sunny	Mild	Normal	Strong	Yes
4	Cloudy	Mild	High	Strong	Yes
5	Rainy	Mild	High	Strong	No
6	Rainy	Cool	Normal	Strong	No
7	Rainy	Mild	High	Weak	Yes
8	Sunny	Hot	High	Strong	No
9	Cloudy	Hot	Normal	Weak	Yes
10	Rainy	Mild	High	Strong	No


Example of Decision Tree

In the below diagram the tree will first ask what is the weather? Is it sunny, cloudy, or rainy? If yes then it will go to the next feature which is humidity and wind. It will again check if there is a strong wind or weak, if it's a weak wind and it's rainy then the person may go and play.



Example of Decision Tree (uncertainty in the dataset)

Now you must be thinking how do I know what should be the root node? what should be the decision node? when should I stop splitting? To decide this, there is a metric called “Entropy” which is the amount of uncertainty in the dataset.



Entropy

Entropy is nothing but the uncertainty in our dataset or measure of disorder.

Suppose you have a group of friends who decides which movie they can watch together on Sunday. There are 2 choices for movies, one is “***Lucy***” and the second is “***Titanic***” and now everyone has to tell their choice. After everyone gives their answer we see that “*Lucy*” gets 4 votes and “*Titanic*” gets 4 votes. Which movie do we watch now?


The formula for Entropy

Here p_+ is the probability of positive class
 p_- is the probability of negative class
 S is the subset of the training example

$$E(S) = -p_{(+)} \log p_{(+)} - p_{(-)} \log p_{(-)}$$

How do Decision Trees use Entropy?

Entropy basically measures the impurity of a node. Impurity is the degree of randomness; it tells how random our data is. A **pure sub-split** means that either you should be getting “yes”, or you should be getting “no”.



How do Decision Trees
use Entropy?...
Is our splits pure?
(For feature 2)

We see here the
split is not pure,
why? Because we
can still see some
negative classes in
both the nodes.

$$\Rightarrow -\left(\frac{5}{7}\right)\log_2\left(\frac{5}{7}\right) - \left(\frac{2}{7}\right)\log_2\left(\frac{2}{7}\right)$$

$$\Rightarrow -(0.71 * -0.49) - (0.28 * -1.83)$$

$$\Rightarrow -(-0.34) - (-0.51)$$

$$\Rightarrow 0.34 + 0.51$$

$$\Rightarrow 0.85$$

How do Decision Trees
use Entropy?...
Is our splits pure?
(For feature 3)

We see here the
split is not pure,
why? Because we
can still see some
negative classes in
both the nodes.

$$\Rightarrow -\left(\frac{3}{5}\right)\log_2\left(\frac{3}{5}\right) - \left(\frac{2}{5}\right)\log_2\left(\frac{2}{5}\right)$$

$$\Rightarrow -(0.6 * -0.73) - (0.4 * -1.32)$$

$$\Rightarrow -(-0.438) - (-0.528)$$

$$\Rightarrow 0.438 + 0.528$$

$$\Rightarrow 0.966$$

Decrease the uncertainty or impurity in the dataset?

the goal of machine learning is to decrease the uncertainty or impurity in the dataset, here by using the entropy we are getting the impurity of a particular node, we don't know if the parent entropy or the entropy of a particular node has decreased or not

For this, we bring a new metric called “Information gain” which tells us how much the parent entropy has decreased after splitting it with some feature.

Information Gain


Information gain measures the reduction of uncertainty given some feature and it is also a deciding factor for which attribute should be selected as a decision node or root node.

It is just entropy of the full dataset – entropy of the dataset given some feature.

$$\text{Information Gain} = E(Y) - E(Y|X)$$


Information Gain Example

suppose our entire population has a total of 30 instances. The dataset is to predict whether the person will go to the gym or not. Let's say 16 people go to the gym and 14 people don't.



Information Gain

Example ...



$$E(\text{Parent}) = -\left(\frac{16}{30}\right)\log_2\left(\frac{16}{30}\right) - \left(\frac{14}{30}\right)\log_2\left(\frac{14}{30}\right) \approx 0.99$$

$$E(\text{Parent}|\text{Energy} = \text{"high"}) = -\left(\frac{12}{13}\right)\log_2\left(\frac{12}{13}\right) - \left(\frac{1}{13}\right)\log_2\left(\frac{1}{13}\right) \approx 0.39$$

$$E(\text{Parent}|\text{Energy} = \text{"low"}) = -\left(\frac{4}{17}\right)\log_2\left(\frac{4}{17}\right) - \left(\frac{13}{17}\right)\log_2\left(\frac{13}{17}\right) \approx 0.79$$

Information Gain

Example ...

To see the weighted average of entropy of each node we will do as follows:

$$E(\textit{Parent}|\textit{Energy}) = \frac{13}{30} * 0.39 + \frac{17}{30} * 0.79 = 0.62$$

Information Gain

Example ...


Now we have the value of $E(\text{Parent})$ and $E(\text{Parent}|\text{Energy})$, information gain will be:

$$\begin{aligned}\text{Information Gain} &= E(\text{parent}) - E(\text{parent}|\text{energy}) \\ &= 0.99 - 0.62 \\ &= 0.37\end{aligned}$$

Our parent entropy was near 0.99 and after looking at this value of information gain, we can say that the entropy of the dataset will decrease by 0.37 if we make “Energy” as our root node.


Information Gain Example-2

Similarly, we will do this with the other feature
“Motivation” and calculate its information gain.



Information Gain

Example 2...



$$E(\text{Parent}) = 0.99$$

$$E(\text{Parent} | \text{Motivation} = \text{"No motivation"}) = -\left(\frac{7}{8}\right)\log_2\left(\frac{7}{8}\right) - \frac{1}{8}\log_2\left(\frac{1}{8}\right) = 0.54$$

$$E(\text{Parent} | \text{Motivation} = \text{"Neutral"}) = -\left(\frac{4}{10}\right)\log_2\left(\frac{4}{10}\right) - \left(\frac{6}{10}\right)\log_2\left(\frac{6}{10}\right) = 0.97$$

$$E(\text{Parent} | \text{Motivation} = \text{"Highly motivated"}) = -\left(\frac{5}{12}\right)\log_2\left(\frac{5}{12}\right) - \left(\frac{7}{12}\right)\log_2\left(\frac{7}{12}\right) = 0.98$$

Information Gain

Example 2...

To see the weighted average of entropy of each node we will do as follows:

$$E(\text{Parent}|\text{Motivation}) = \frac{8}{30} * 0.54 + \frac{10}{30} * 0.97 + \frac{12}{30} * 0.98 = 0.86$$

Information Gain

Example 2...

Now we have the value of $E(\text{Parent})$ and $E(\text{Parent}|\text{Energy})$, information gain will be:

$$\begin{aligned}\text{Information Gain} &= E(\text{Parent}) - E(\text{Parent}|\text{Motivation}) \\ &= 0.99 - 0.86 \\ &= 0.13\end{aligned}$$

Our parent entropy was near 0.99 and after looking at this value of information gain, we can say that the entropy of the dataset will decrease by 0.37 if we make “Energy” as our root node.

Conclusion based on Information Gain in gym Example

We now see that the “Energy” feature gives more reduction which is 0.37 than the “Motivation” feature. Hence we will select the feature which has the highest information gain and then split the node based on that feature.

In this example “Energy” will be our root node and we’ll do the same for sub-nodes. Here we can see that when the energy is “high” the entropy is low and hence we can say a person will definitely go to the gym if he has high energy, but what if the energy is low? We will again split the node based on the new feature which is “Motivation”.

6

Logistic Regression

7

Decision Tree

8

Support Vector Machine (SVM)

9


Data Preparation & Metric Trap

10

Training & Evaluation

SVM


Support Vector Machine (SVM) is a model generally used for **classification and regression problems**. It can solve **linear and nonlinear** problems that help the algorithm to tackle real-time projects.



SVM

How It's works?
(linear models)


we will plot all the data points with class, and then we draw multiple hyperplanes, here, it will try to separate the two classes.



SVM

How It's works? (non-linear models)

We can see that the data point and class won't be separated by a hyperplane. We have multiple techniques to handle nonlinear decision boundaries using kernels.



Basic mathematics behind the Linear Hard-Margin Classifier(Linear SVM) for Strictly Separable Case

$$d(x, w, b) = w^T x + b = \sum_{i=1}^n w_i x_i + b$$

Where, $d(x, w, b)$ is the distance of point x from the hyperplane.
Greater the distance stronger the classification of x .

1. $d(x, w, b) > 0$, classify x as class 1 (i.e. its associated $y = +1$)
2. $d(x, w, b) < 0$, classify x as class 2 (i.e. its associated $y = -1$)
3. $d(x, w, b) = 0$, then x is equi-probable for both classes so, as per the business problem we classify this data point

6

Logistic Regression

7

Decision Tree

8

Support Vector Machine (SVM)

9

Data Preparation & Metric Trap

10

Training & Evaluation

Data Preparation

We already have decided with the list of columns we will use to for training as shown below:

```
col_set1 = ["robust_amount", "robust_time", "V2",  
            "V4", "V10", "V11", "V12", "V14", "V16", "V17", "V19"]  
col_set2 = ["ptrans_amount", "ptrans_time", "V2",  
            "V4", "V10", "V11", "V12", "V14", "V16", "V17", "V19"]
```

SMOTE data preparation

We will have two sets of columns, and we will also use different resampled data frame. In this case, we will use SMOTE and Random under-sampled dataset.

```
X_train_smote_robust = train_df_sm[col_set1]  
y_train_smote_robust = train_df_sm["Class"]  
  
X_train_smote_power = train_df_sm[col_set2]  
y_train_smote_power = train_df_sm["Class"]
```

Random under-sample data preparation

We will have two sets of columns, and we will also use different resampled data frame. In this case, we will use SMOTE and Random under-sampled dataset.

```
X_train_under_robust = train_df_under[col_set1]  
y_train_under_robust = train_df_under["Class"]  
  
X_train_under_power = train_df_under[col_set2]  
y_train_under_power = train_df_under["Class"]
```


Metric Trap

Here “accuracy” cannot be used as a metric because this is an imbalanced dataset.

```
def metric_trap_acc(x=None):  
    return 0 #Class 0
```

Pseudo model (should not be used)

The above function/pseudo-model has great accuracy as it will always return class zero. The probability of getting class zero for the imbalanced data set is more than 95%.

6

Logistic Regression

7

Decision Tree

8

Support Vector Machine (SVM)

9

Data Preparation & Metric Trap

10

Training & Evaluation

Preparation of training data

We will train all the models with four different kinds of the dataset and evaluate some of the models

```
data = [  
    [X_train_smote_robust, y_train_smote_robust, "SMOTE -Robust Scaling"],  
    [X_train_smote_power, y_train_smote_power, "SMOTE -Power Transformer"],  
    [X_train_under_robust, y_train_under_robust, "Under Sampling -Robust Scaling"],  
    [X_train_under_power, y_train_under_power, "Under Sampling -Power Transformer"]  
]
```

Initializing models

We are planning to write a small code that will help us to train all the models for all the data at once using cross validation

```
# Classifier Libraries
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier

classifiers = {
    "DecisionTreeClassifier": DecisionTreeClassifier(),
    "LogisiticRegression": LogisticRegression(max_iter=1000),
    "KNearest": KNeighborsClassifier(),
    "Radial Basis Support Vector Classifier": SVC()
}
```

Fitting and cross-validating

We have a separate set of testing data, but we need to test with the training data as well, using the k-fold cross-validation technique. This will decrease the chance of selection bias. With this process, we will select the optimal dataset and some of the models for the evaluation.

```
from sklearn.model_selection import cross_val_score

for X, y, name in data:
    print("\n\n" + name + ":\n")
    for key, classifier in classifiers.items():
        classifier.fit(X, y)
        training_score = cross_val_score(classifier, X, y, cv=5)
        print("Classifiers: ", classifier.__class__.__name__, " has a training score of", \
              round(training_score.mean(), 2) * 100, "% accuracy score")
```

Cross-validation training score for SMOTE

We can see that k-NN's accuracy is 100%, and it is highly probable and possible that it is suffering from overfitting.

SMOTE -Robust Scaling:

```
Classifiers: DecisionTreeClassifier has a training score of 80.0 % accuracy score
Classifiers: LogisticRegression has a training score of 97.0 % accuracy score
Classifiers: KNeighborsClassifier has a training score of 100.0 % accuracy score
Classifiers: SVC has a training score of 98.0% accuracy score
```

SMOTE -Power Transformer:

```
Classifiers: DecisionTreeClassifier has a training score of 80.0 % accuracy score
Classifiers: LogisticRegression has a training score of 97.0 % accuracy score
Classifiers: KNeighborsClassifier has a training score of 100.0 % accuracy score
Classifiers: SVC has a training score of 98.0% accuracy score
```


Cross-validation training score for under-sampling

We can use Decision Tree, Logistic Regression for SMOTE, and Support Vector Machine, k-NN for Random Under Sampling

Under Sampling -Robust Scaling:

```
Classifiers: DecisionTreeClassifier has a training score of 91.0 % accuracy score
Classifiers: LogisticRegression has a training score of 93.0 % accuracy score
Classifiers: KNeighborsClassifier has a training score of 94.0 % accuracy score
Classifiers: SVC has a training score of 94.0% accuracy score
```

Under Sampling -Power Transformer:

```
Classifiers: DecisionTreeClassifier has a training score of 90.0 % accuracy score
Classifiers: LogisticRegression has a training score of 94.0 % accuracy score
Classifiers: KNeighborsClassifier has a training score of 94.0 % accuracy score
Classifiers: SVC has a training score of 94.0% accuracy score
```


Training models

The order of the features in the dataset during training and testing should be the same else it will give us a wrong result

```
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier

SVC_Under_Sampling_Power = SVC(shrinking=False, probability=True)
kNN_Under_Sampling_Power = KNeighborsClassifier()

DT_SMOTE_Robust = DecisionTreeClassifier()
LR_SMOTE_Robust = LogisticRegression(max_iter=1000)

SVC_Under_Sampling_Power.fit(X_train_under_power, y_train_under_power)
kNN_Under_Sampling_Power.fit(X_train_under_power, y_train_under_power)

DT_SMOTE_Robust.fit(X_train_smote_robust, y_train_smote_robust)
LR_SMOTE_Robust.fit(X_train_smote_robust, y_train_smote_robust)
```

Random under-sample model configuration

SVC_Under_Sampling_Power

```
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,  
    decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',  
    max_iter=-1, probability=True, random_state=None, shrinking=False,  
    tol=0.001, verbose=False)
```

kNN_Under_Sampling_Power

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',  
                     metric_params=None, n_jobs=None, n_neighbors=5, p=2,  
                     weights='uniform')
```

SMOTE Model configuration

DT_SMOTE_Robust

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',  
                      max_depth=None, max_features=None, max_leaf_nodes=None,  
                      min_impurity_decrease=0.0, min_impurity_split=None,  
                      min_samples_leaf=1, min_samples_split=2,  
                      min_weight_fraction_leaf=0.0, presort='deprecated',  
                      random_state=None, splitter='best')
```

LR_SMOTE_Robust

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
                  intercept_scaling=1, l1_ratio=None, max_iter=1000,  
                  multi_class='auto', n_jobs=None, penalty='l2',  
                  random_state=None, solver='lbfgs', tol=0.0001, verbose=0,  
                  warm_start=False)
```

Receiver Operating Characteristics(ROC)

We had discussed before the metric trap, and just using **accuracy will give us wrong information**. There are multiple metrics we can see. But we will use Receiver Operating Characteristics (ROC) to analyze all the models.

ROC curve is a plot between True Positive Rate and False Positive Rate at various threshold values

ROC is a way to analyze the performance of a binary classifier system as its discrimination threshold is valid.

Generating values for plotting ROC curve

The below code snippet will **give true positive rates, false-positive rates, and the threshold for all the models**. Now we have a task to plot the data and visualize it.

```
from sklearn.metrics import roc_curve

log_fpr, log_tpr, log_threshold = roc_curve(y_test,
                                            LR_SMOTE_Robust.predict_proba(X_test_robust)[:,-1])
knear_fpr, knear_tpr, knear_threshold = roc_curve(y_test,
                                                  KNN_Under_Sampling_Power.predict_proba(X_test_power)[:,-1])
svc_fpr, svc_tpr, svc_threshold = roc_curve(y_test,
                                            SVC_Under_Sampling_Power.predict_proba(X_test_power)[:,-1])
tree_fpr, tree_tpr, tree_threshold = roc_curve(y_test,
                                                DT_SMOTE_Robust.predict_proba(X_test_robust)[:,-1])
```



Plotting ROC curve

We will only plot the ROC curve, and that is good for analyzing the model, but it is recommended to plot some visualization with a threshold as one axis.

```
plt.figure(figsize=(16,8))
plt.title('ROC Curve \n Top 4 Classifiers', fontsize=18)
plt.plot(log_fpr, log_tpr,
         label='LR_SMOTE_Robust: {:.4f}'.format( \
             roc_auc_score(y_test, LR_SMOTE_Robust.predict_proba(X_test_robust)[: ,1])))
plt.plot(tree_fpr, tree_tpr,
         label='DT_SMOTE_Robust: {:.4f}'.format( \
             roc_auc_score(y_test, DT_SMOTE_Robust.predict_proba(X_test_robust)[: ,1])))
plt.plot(knear_fpr, knear_tpr,
         label='kNN_Under_Sampling_Power: {:.4f}'.format( \
             roc_auc_score(y_test, kNN_Under_Sampling_Power.predict_proba(X_test_power)[: ,1])))
plt.plot(svc_fpr, svc_tpr,
         label='SVC_Under_Sampling_Power: {:.4f}'.format( \
             roc_auc_score(y_test, SVC_Under_Sampling_Power.predict_proba(X_test_power)[: ,1])))
plt.plot([0, 1], [0, 1], 'k--')
plt.axis([-0.01, 1, 0, 1])
plt.xlabel('False Positive Rate', fontsize=16)
plt.ylabel('True Positive Rate', fontsize=16)
plt.annotate('Minimum ROC Score of 50%', xy=(0.5, 0.5), xytext=(0.6, 0.3),
            arrowprops=dict(facecolor='#6E726D', shrink=0.05),)
plt.legend()
plt.show()
```

ROC curve (AUC or AUROC)

The higher the value of ROC, the better the model is. We have drawn a reference line with 0.5 as a ROC score. It means that lines/models close to the reference line signifies a random guess



Generating AUROC Score

ROC score is nothing but the area under the curve, which is generated from the ROC curve.

```
from sklearn.metrics import roc_auc_score

print("ROC Accuracy Score:\n")

print('LR_SMOTE_Robust: ', roc_auc_score(y_test,
                                          LR_SMOTE_Robust.predict_proba(X_test_robust)[: ,1]))
print('DT_SMOTE_Robust: ', roc_auc_score(y_test,
                                          DT_SMOTE_Robust.predict_proba(X_test_robust)[: ,1]))
print('KNN_Under_Sampling_Power: ', roc_auc_score(y_test,
                                                    KNN_Under_Sampling_Power.predict_proba(X_test_power)[: ,1]))
print('SVC_Under_Sampling_Power: ', roc_auc_score(y_test,
                                                    SVC_Under_Sampling_Power.predict_proba(X_test_power)[: ,1]))
```

ROC Score


We can see from the score and curve that Decision Tree performed worst hence we can drop that model. The best models as per score is SVC and Logistic Regression

```
LR_SMOTE_Robust: 0.9672841556352365
DT_SMOTE_Robust: 0.5421638247412819
kNN_Under_Sampling_Power: 0.9328843189132074
SVC_Under_Sampling_Power: 0.9770356462621923
```

Precision-Recall curve

ROC curves are generally used when there is a **roughly equal number of records** for each class.


Precision-Recall curves are used when there is a **moderate to high-class imbalance**.



Confusion Matrix (Decision Tree + SMOTE + Robust Scaling)

Reading the matrix, we can mostly say all values of Class 1 are misclassified, and the model is extremely biased towards Class 0. Seeing the performance, we will surely drop this model and not consider it for further re-tuning.


```
from sklearn.metrics import plot_confusion_matrix  
plot_confusion_matrix(DT_SMOTE_Robust, X_test_robust, y_test, cmap=plt.cm.Blues)  
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a6ce86d50>
```



Confusion Matrix (k-NN + Random Under Sampling + Power Transformer)

For k-NN, the misclassification for Class 0 is a bit high compared to Logistic and SVC. So, for this reason, we will be rejecting this model too

```
plot_confusion_matrix(kNN_Under_Sampling_Power, X_test_power, y_test, cmap=plt.cm.Blues)  
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a6bed7710>
```




SVC + Random Under Sampling + Power Transformer

SVC is similar to Logistic Regression, and we see the difference is minimal, and both can be used interchangeably.

```
plot_confusion_matrix(SVC_Under_Sampling_Power, X_test_power, y_test, cmap=plt.cm.Blues)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a70733e50>
```



Generating a classification report

Now, we need to choose between SVC and Logistic Regression. There is another way we can get all the important metrics for analysis. We can use the Classification Report

```
from sklearn.metrics import classification_report

print('LR_SMOTE_Robust:')
print(classification_report(y_test, LR_SMOTE_Robust.predict(X_test_robust)))

print('DT_SMOTE_Robust:')
print(classification_report(y_test, DT_SMOTE_Robust.predict(X_test_robust)))

print('SVC_Under_Sampling_Power:')
print(classification_report(y_test, SVC_Under_Sampling_Power.predict(X_test_power)))

print('kNN_Under_Sampling_Power:')
print(classification_report(y_test, kNN_Under_Sampling_Power.predict(X_test_power)))
```


SMOTE classification report

The **recall** means "how many of this class you find over the whole number of element of this class"

The **precision** will be "how many are correctly classified among that class"

The **f1-score** is the **harmonic mean between precision & recall**

The **support** is the number of occurrence of the given class in your dataset

LR_SMOTE_Robust:				
	precision	recall	f1-score	support
0	1.00	0.99	1.00	28251
1	0.19	0.79	0.30	47
accuracy			0.99	28298
macro avg	0.59	0.89	0.65	28298
weighted avg	1.00	0.99	1.00	28298
DT_SMOTE_Robust:				
	precision	recall	f1-score	support
0	1.00	1.00	1.00	28251
1	0.15	0.09	0.11	47
accuracy			1.00	28298
macro avg	0.58	0.54	0.55	28298
weighted avg	1.00	1.00	1.00	28298

Random under-sampling classification report

The **recall** means "how many of this class you find over the whole number of element of this class"

The **precision** will be "how many are correctly classified among that class"

The **f1-score** is the **harmonic mean between precision & recall**

The **support** is the number of occurrence of the given class in your dataset

```
SVC_Under_Sampling_Power:
      precision    recall  f1-score   support


     0       1.00      0.99      1.00     28251
     1       0.20      0.79      0.32         47

 accuracy          0.99     28298
 macro avg          0.60      0.89      0.66     28298
weighted avg          1.00      0.99      1.00     28298

kNN_Under_Sampling_Power:
      precision    recall  f1-score   support

     0       1.00      0.99      1.00     28251
     1       0.13      0.83      0.23         47

 accuracy          0.99     28298
 macro avg          0.57      0.91      0.61     28298
weighted avg          1.00      0.99      0.99     28298
```



Appendix A: Cross Validation and K-fold Method

Cross-validation

Cross-validation is a technique by which we can assess how a model will perform in practice, i.e., in real-time and unseen data

Cross-validation is a model validation technique to ensure its performance, but it only works best when the training dataset distribution is similar to real-time data distribution.


There are different types of cross-validation techniques:

Exhaustive cross-validation like Leave-p-out cross-validation

Non-exhaustive cross-validation like k-fold and Holdout method

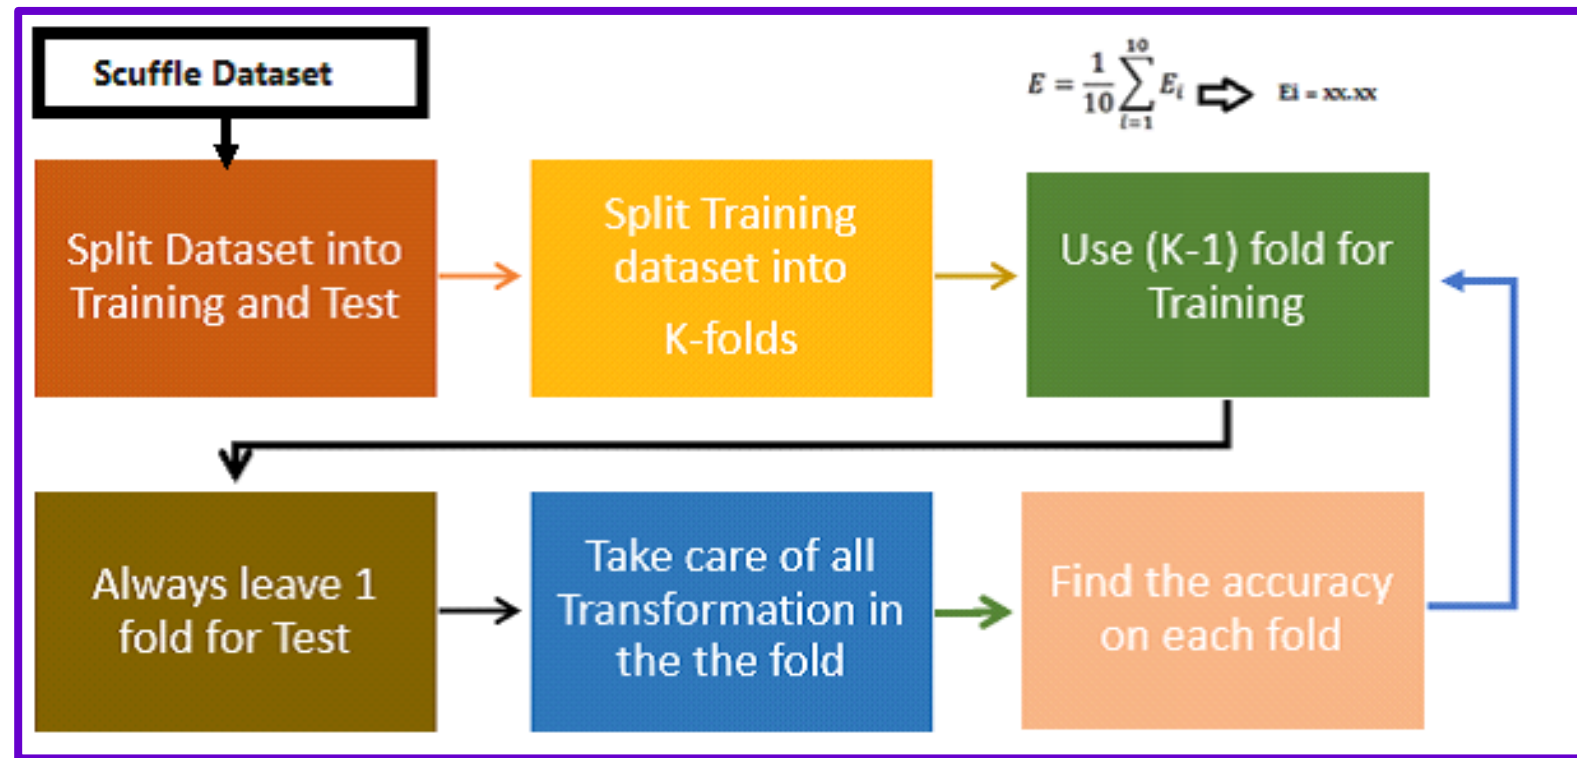
k-fold cross-validation

K-fold cross-validation is a technique for evaluating predictive models. The dataset is divided into k subsets or folds.



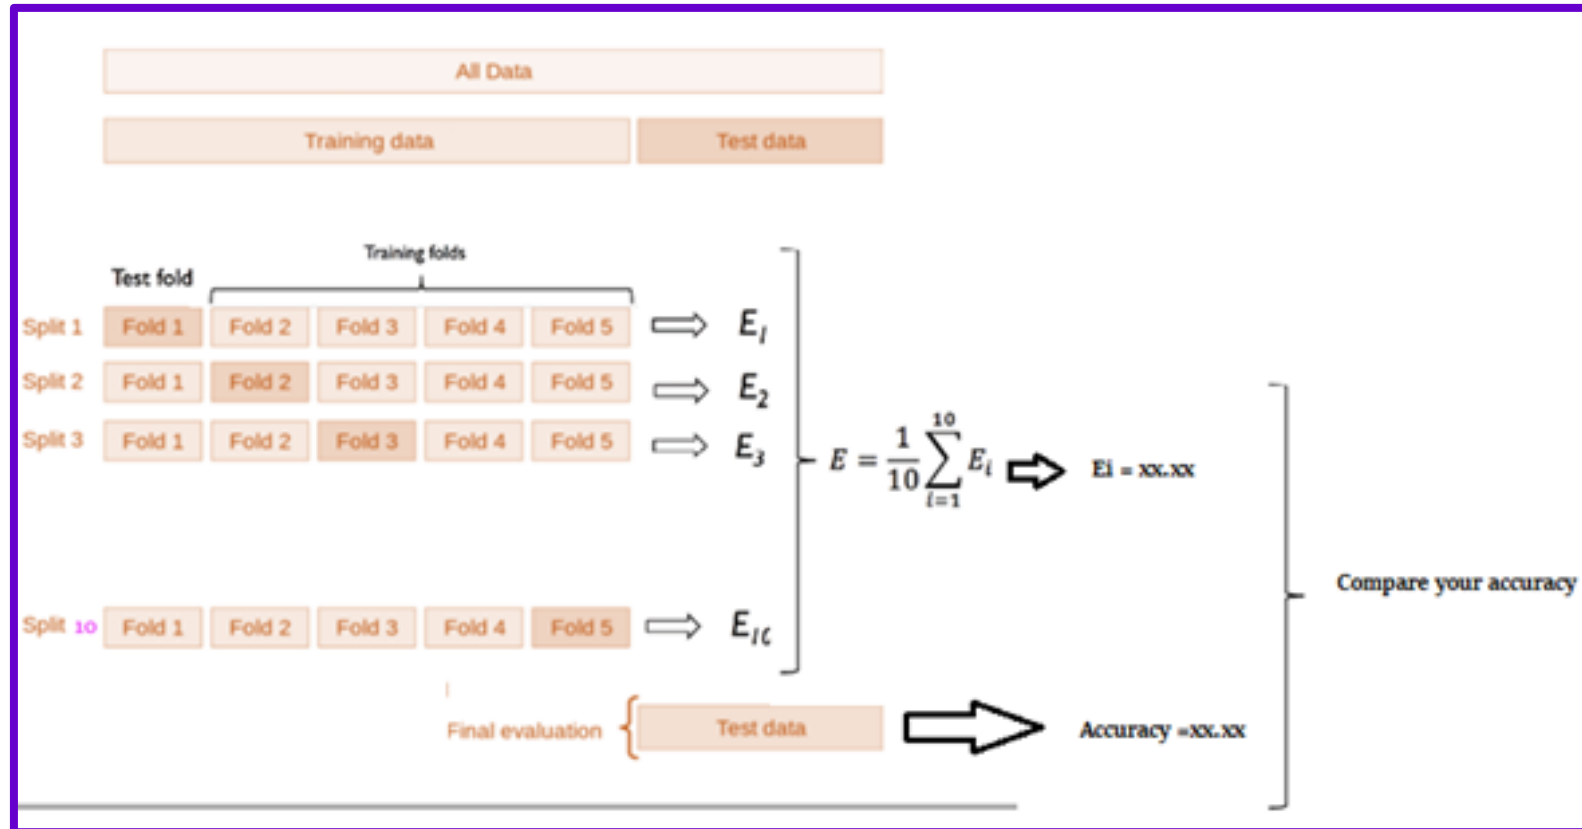
k-fold cross-validation

The model is trained and evaluated k times, using a different fold as the validation set each time.




k-fold cross-validation

The model is trained and evaluated k times, using a different fold as the validation set each time.



k-fold cross-validation Example

If $K=5$, it means, in the given dataset and we are splitting into 5 folds and running the Train and Test. During each run, one fold is considered for testing and the rest will be for training



Designed by Author - Shanthababu

Course References

- [1] S. J. Russell and P. Norvig, *Artificial Intelligence: A Modern Approach*. Pearson, 2021.
- [2] T. Ghosh and S. K. B. Math, *Practical Mathematics for AI and Deep Learning: A Concise yet In-Depth Guide on Fundamentals of Computer Vision, NLP, Complex Deep Neural Networks and Machine Learning (English Edition)*. BPB Publications, 2022.
- [3] M. P. Deisenroth, A. A. Faisal, and C. S. Ong, *Mathematics for Machine Learning*. Cambridge University Press, 2020.
- [4] T. V. Geetha and S. Sendhilkumar, *Machine Learning: Concepts, Techniques and Applications*. CRC Press LLC, 2023.
- [5] A. Géron, *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems*. O'Reilly Media, 2023.
- [6] O. Theobald, *Machine Learning for Absolute Beginners: A Plain English Introduction (Third Edition)*. Scatterplot Press, 2021.

Accessing Course Resource



[linkedin.com/in/Samanipour](https://www.linkedin.com/in/Samanipour)



t.me/SamaniGroup



github.com/Samanipour